

Back to Basics: Modern C++ Style

loops
pointers & references
smart pointers
variable declarations
parameter passing

Herb Sutter

CA

Complexity Anonymous

A 12-step program
for good people attempting to
recover from complexity addiction



It's hard to ~~be~~ remember you're an expert

C++ developers (~3M)

libstdc++ developers (~30)
+
libc++ developers (~5-7)
+
Boost developers (~300?)
+
ISO WG21 attendees (~300?)

Reality Check




Occurrences of “&&” in Bjarne’s 90-min Tue keynote? **0**

Value of modern C++’s simple usable defaults? **Priceless**

Most Important C++ Book

- ▶ “What should every C++ programmer be expected to know?”
 - ▶ For years, there has not been a single source to point to.
- ▶ Now there is. In 180 pages you can read on a long plane flight.
 - ▶ **Recommend it heavily!**
 - ▶ Also a demonstration that modern C++ is simpler to teach and explain.



This Talk

- ▶ This talk focuses on **defaults**, basic styles and idioms in modern C++.
 - ✗ ~~“Default” != “don’t think.”~~
 - ✓ “Default” == “don’t **overthink.**” Esp. don’t **optimize prematurely.**
- ▶ These reinforce (not compete with) the “fundamentals.”
 - ▶ “Write for clarity and correctness first.”
 - ▶ “Avoid premature optimization.” By default, prefer *clear* over *optimal*.
 - ▶ “Avoid premature pessimization.” Prefer *faster* when equally clear.

Prefer range-for

why do this
`for(auto i = begin(c); i != end(c); ++i) { ... use(*i); ... }`

when you can do this
`for(auto& e : c) { ... use(e); ... }`

and soon this
`for(e : c) { ... use(e); ... }`

Use smart pointers effectively...

... but still **use lots of raw * and &**, they're great!

wait, what?

Don't Use Owning *, *new*, or *delete*

▶ C++98:

```
widget* factory();  
void caller() {  
    widget* w = factory();  
    gadget* g = new gadget();  
    use( *w, *g );  
    delete g;  
    delete w;  
}
```

red ⇒ now “mostly wrong” 😊

▶ Don't use owning *,
new or *delete*.

- ▶ Except: Encapsulated inside the implementation of low-level data structures.

▶ Modern C++:

```
unique_ptr<widget> factory();  
void caller() {  
    auto w = factory();  
    auto g = make_unique<gadget>();  
    use( *w, *g );  
}
```

▶ For “new”, use *make_unique* by default,
make_shared if it will be shared.

- ▶ For “delete”, write nothing.

Don't Use Owing ***, *new*, or *delete*

NB: important qualifier

▶ C++98:

```
widget* factory();
void caller() {
    widget* w = factory();
    gadget* g = new gadget();
    use( *w, *g );
    delete g;
    delete w;
}
red ⇒ now “mostly wrong” ☺
```

▶ Don't use owning ***, *new* or *delete*.

- ▶ Except: Encapsulated inside the implementation of low-level data structures.

▶ C++14:

```
unique_ptr<widget> factory();
void caller() {
    auto w = factory();
    auto g = make_unique<gadget>();
    use( *w, *g );
}
```

▶ For “new”, use *make_unique* by default, *make_shared* if it will be shared .

- ▶ For “delete”, write nothing.

NB: Non-Owning ***/*&* Are Still Great

▶ C++98 “Classic”:

```
void f( widget& w ) { // if required
    use(w);
}
void g( widget* w ) { // if optional
    if(w) use(*w);
}
```



*** and *&* FTW

(More on parameter passing coming later...)

▶ Modern C++ “Still Classic”:

```
void f( widget& w ) { // if required
    use(w);
}
void g( widget* w ) { // if optional
    if(w) use(*w);
}
```

```
auto upw = make_unique<widget>();
...
f( *upw );

auto spw = make_shared<widget>();
...
g( spw.get() );
```

Antipatterns Hurt Pain Pain

- ▶ Antipattern #1: Parameters
(Note: *Any* refcounted pointer type.)

```
void f( refcnt_ptr<widget>& w ) {  
    use(*w);  
} // ?  
void f( refcnt_ptr<widget> w ) {  
    use(*w);  
} // ?!?!
```
- ▶ Antipattern #2: Loops
(Note: *Any* refcounted pointer type.)

```
refcnt_ptr<widget> w = ...;  
for( auto& e: baz ) {  
    auto w2 = w;  
    use(w2,*w2,w,*w,whatever);  
} // ?!?!?!?
```

Example (thanks Andrei): In late 2013, Facebook RocksDB changed from pass-by-value *shared_ptr* to pass-*/&. QPS improved 4× (100K to 400K) in one benchmark.

<http://tinyurl.com/gotw91-example>

No Copy No Cry

FAQ: Smart Pointer Parameters — See GotW #91 (tinyurl.com/gotw91)

RefCounted smart pointers are about managing the **owned object's lifetime**. Copy/assign one only when you intend to manipulate the **owned object's lifetime**.

Any “smart pointers (or `std::vectors`) are slow” performance claims based on code that copies/assigns smart pointers (or `std::vectors`) – including **passing by value** or **copying/assigning in loops** – when copies are not needed are fundamentally flawed.

Yes, this applies to your refcounted smart pointer:

- `shared_ptr` (Boost, TR1, `std::`)
- `retain/release` (Objective-C ARC, Clang 3.5)
- `AddRef/Release` (COM and WinRT, C++/CX ^)
- any other refcounting strategy you will ever see

Passing Smart Pointers (G++ W #91)

```

unique_ptr<widget> factory(); // source – produces widget
void sink( unique_ptr<widget> ), // sink – consumes widget
void reseal( unique_ptr<widget>& ); // “will” or “might” reseal ptr
void thinko( const unique_ptr<widget>& ); // usually not what you want

shared_ptr<widget> factory(); // source + shared ownership
// when you know it will be shared, perhaps by factory itself
void share( shared_ptr<widget> ); // share – “will” retain refcount
void reseal( shared_ptr<widget>& ); // “will” or “might” reseal ptr
void may_share( const shared_ptr<widget>& ); // “might” retain refcount

```

&&? wait for it...

How to “Do It Right” (Partial)

1. Never pass smart pointers (by value or by reference) unless you actually want to manage their lifetime, store, change, or let go of a reference.
 - ▶ Prefer pointer-to-const
 - ▶ Else if you need to modify, use pointer-to-non-const
2. Express ownership semantics when you don't want to manage their lifetime.
 - ▶ It's free
 - ▶ It's safe
 - ▶ It's declarative = expresses intent
 - ▶ It removes many (often most) objects out of the programmer's mental model.
3. Else use `make_shared` up front wherever possible, if object will be shared.

Not quite done: One guideline missing...
...and it applies to any RC pointer type,
in almost any language / library

Guideline: Dereference *Unaliased+Local* RC Ptrs

- ▶ The reentrancy pitfall (simplified):
- ▶ “Pin” using unaliased local copy.

```
// global (static or heap), or aliased local
... shared_ptr<widget> g_p ...
```

```
// global (static or heap), or aliased local
... shared_ptr<widget> g_p ...
```

```
void f( widget& w ) {
    g();
    use(w);
}
void g() {
    g_p = ...;
}
```

```
void f( widget& w ) {
    g();
    use(w);
}
void g() {
    g_p = ...;
}
```

```
void my_code() {
    f( *g_p );    // passing *nonlocal
}
// should not pass code review
```

```
void my_code() {
    auto pin = g_p; // 1 ++ for whole tree
    f( *pin );     // ok, *local
}

```

Guideline: Dereference *Unaliased+Local* RC Ptrs

- ▶ The reentrancy pitfall (simplified):
- ▶ “Pin” using unaliased local copy.

```
// global (static or heap), or aliased local
... shared_ptr<widget> g_p ...
```

```
// global (static or heap), or aliased local
... shared_ptr<widget> g_p ...
```

```
void f( widget& w ) {
    g();
    use(w);
}
void g() {
    g_p = ...;
}
```

```
void f( widget& w ) {
    g();
    use(w);
}
void g() {
    g_p = ...;
}
```

```
void my_code() {
    f( *g_p );    // passing *nonlocal
    g_p->foo();   // (or nonlocal->)
}
// should not pass code review
```

```
void my_code() {
    auto pin = g_p; // 1 ++ for whole tree
    f( *pin );     // ok, *local
    pin->foo();    // ok, local->
}

```

Summary: How to “Do It Right”


1. **Never pass smart pointers (by value or by reference) unless you actually want to manipulate the pointer ⇒ store, change, or let go of a reference.**
 - ▶ Prefer passing objects by * or & as usual – just like always.
*Remember: Take **unaliaised+local copy** at the top of a call tree, don't pass f(*g_p).*
 - ▶ **Else** if you do want to manipulate lifetime, great, do it as on previous slide.
2. **Express ownership using `unique_ptr` wherever possible**, including when you don't know whether the object will actually ever be shared.
 - ▶ It's free = exactly the cost of a raw pointer, by design.
 - ▶ It's safe = better than a raw pointer, including exception-safe.
 - ▶ It's declarative = expresses intended uniqueness and source/sink semantics.
 - ▶ **It removes many (often most) objects out of the ref counted population.**
3. **Else use `make_shared` up front wherever possible**, if object will be shared.

Write **`make_unique`** (by default)
or **`make_shared`** (when needed)
instead of *new* and *delete*.

Don't use owning raw *, *new*, or *delete* any more, except rarely inside the implementation details of low-level data structures.

Do use non-owning raw * and &, especially for parameters.

Don't copy/assign refcounted smart pointers, including pass-by-value or in loops, unless you really want the semantics they express: altering object lifetime.



Let's talk about *auto*...




It's okay, it's *really simple*...

Spoiler

To make type track, deduce:
***auto** var = init;*

To make type stick, commit:
*auto var = **type**{ init };*
or ***type** var{ init };*



Consider This Code

- ▶ **Guru Meditation Q:** What does this code do?

```
template<class Container, class Value>  
void ██████████( Container& c, Value v )  
{  
    if( find(begin(c), end(c), v) == end(c) )  
        c.push_back( move(v) );  
    assert( !c.empty() );  
}
```

Why Not “Just Deduce the Type”?

▶ Counterarguments: “Oi, but it’s unreadable!” “What’s my type?”

▶ This is a weak argument for three reasons:

- ▶ (Minor) It doesn’t matter to anyone who uses an IDE.
- ▶ (Major) It reflects bias to code against **implementations**, not **interfaces**.
- ▶ (Major) We already ignore actual types with **templates** and **temporaries**.


```
template<class Container, class Value> // what type is Container? Value?
void append_unique( Container& c, Value v ) // anything usable like this...
{
    if( find(begin(c), end(c), v) == end(c) ) // what type does find return?
        c.push_back( move(v) );           // anything comparable to end(cont)...
    assert( !c.empty() ); // what type does .empty() return?
} // anything testable like a bool...

```

 - ▶ We also ignore actual types with virtual functions, *function<>*, etc.

Why Deduce: (1) Correctness

▶ With deduction you always get right type. *Repetition ∝ P(lying)*

▶ Example:

```
void f( const vector<int>& v ) {
    vector<int>::iterator i = v.begin(); // ?
}

```

▶ Options:

```
void f( const vector<int>& v ) {
    vector<int>::iterator i = v.begin(); // error
    vector<int>::const_iterator i = v.begin(); // ok + extra thinking
    auto i = v.begin(); // ok, default
}

```

Why Deduce: (2) Correctness + Maintainability

- ▶ Using deduction makes your code more robust in the face of change.

- ▶ Deduction tracks the correct type when an expression's type changes.
- ▶ Committing to explicit type = silent conversions, needless build breaks.

- ▶ Examples:

```
int i = f(1,2,3) * 42;           // before: ok enough
int i = f(1,2,3) * 42.0;       // after: silent narrowing conversion
auto i = f(1,2,3) * 42.0;      // after: still ok, tracks type

widget w = factory();          // before: ok enough, returns a widget
widget w = factory();          // after: silent conversion, returns a gadget
auto w = factory();            // after: still ok, tracks type

map<string,string>::iterator i = begin(dict); // before: ok enough
map<string,string>::iterator i = begin(dict); // after: error, unordered_map
auto i = begin(dict);          // after: still ok, tracks type
```

Why Deduce: (3) Performance

- ▶ Deduction guarantees no implicit conversion will happen.
 - ▶ A.k.a. “guarantees better performance by default.”
 - ▶ Committing to an explicit type that requires a conversion means silently getting a conversion whether you expected it or not.

Why Deduce: (4) Usability

- ▶ Using deduction is your only good (usable and efficient) option for hard-to-spell and unutterable types like:
 - ▶ lambdas,
 - ▶ binders,
 - ▶ *detail::* helpers,
 - ▶ template helpers, such as expression templates (when they should stay unevaluated for performance), and
 - ▶ template parameter types, which are anonymized anyway,
- ▶ ... short of resorting to:
 - ▶ repetitive *decltype* expressions, and
 - ▶ more-expensive indirections like *std::function*.

Why Deduce: (5) Convenience

- ▶ And, yes, “basic deduction” *auto x = expr;* syntax is almost always less typing.
 - ▶ Mentioned last for completeness because it’s a common reason to like it, but it’s not the biggest reason to use it.

Why Deduce: Wrapup

- ▶ **Prefer *auto x = expr;* by default on variable declarations.**
 - ▶ It offers so much correctness, clarity, maintainability, performance and simplicity goodness that you're only hurting yourself (and your code's future maintainers) if you don't.
 - ▶ Prefer to habitually **program against interfaces, not implementations**. We do this all the time in temporaries and templates anyway and nobody bats an eye.
- ▶ **But: Do commit to an explicit type when you really mean it, which nearly always means you want an explicit conversion.**
 - ▶ **Q:** But even then, does "commit to an explicit type" mean "don't use *auto*"?

Left-to-right auto style

- ▶ Deduce to track if you don't need to commit to a type:
 - ▶ `const char* s = "Hello";` **auto** `s = "Hello";`
 - ▶ `widget w = get_widget();` **auto** `w = get_widget();`
- ▶ Commit to stick to a specific type. Try it **on the right** (same syntax order):
 - ▶ `employee e{ empid }; auto e = employee{ empid };`
 - ▶ `widget w{ 12, 34 }; auto w = widget{ 12, 34 };`
- ▶ With heap allocation, type is **on the right** naturally anyway:
 - ▶ C++98 style: **auto** `w = new widget{};`
 - ▶ C++14 style: **auto** `w = make_unique<widget>();`
- ▶ Teaser: Does this remind you of anything else in C++11? and C++14?
 - ▶ `int f(double); auto f(double) -> int; // C++11`
 - ▶ **auto** `f(double) { ... } // C++14`

The Elephant

But what about

```
int x = 42;
```

vs.

```
auto x = 42;
```

?

“OBVIOUSLY `int x = 42;`
is the tersest and clearest style.”

Right?



Left-to-right auto style

```
employee e{ empid };  
widget w = get_widget();
```

```
auto e = employee{ empid };  
auto w = get_widget();
```

- ▶ Now consider literal suffixes:

```
int x = 42;  
float x = 42.;  
unsigned long x = 42;  
string x = "42";  
chrono::nanoseconds x{ 42 };
```

```
auto x = 42;  
auto x = 42.f;           // no narrowing  
auto x = 42ul;  
auto x = "42"s;         // C++14  
auto x = 42ns;         // C++14
```

- ▶ Remember functions, lambdas, and aliases:

```
int f( double );  
  
typedef set<string> dict;  
template<class T> struct myvec  
{ typedef vector<T,myalloc> type; };
```

```
auto f ( double ) -> int;  
auto f = [=]( double ) { /* ... */ };  
using dict = set<string>;  
template<class T>  
using myvec = vector<T,myalloc>;
```

Left-to-right modern C++ style

- ▶ The C++ world is moving to left-to-right everywhere:

category name = **type** *and/or* initializer ;

Auto variables:

```
auto e = employee{ empid };  
auto w = get_widget();
```

Literals:

```
auto x = 42;  
auto x = 42.f;  
auto x = 42ul;  
auto x = "42"s;  
auto x = 1.2ns;
```

User-defined literals:

Function declarations:

```
auto func ( double ) -> int;
```

Named lambdas:

```
auto func = [=]( double ) { /*...*/ };
```

Aliases (*no more typedefs*):

```
using dict = set<string>;
```

Template aliases:

```
template<class T>  
using myvec = vector<T,myalloc>;
```

I Know Some of You Have Been Wondering

- ▶ Consider:

```
auto x = value;
```

- ▶ Q: Does this "=" create a temporary object plus a move/copy?
 - ▶ Standard says "No." The code $T x = a;$ has exactly the same meaning as $T x(a);$ when a has type T (or derived from T)... and $auto x = a;$ guarantees the types are the same (yay *auto*) so it always means exactly the same as $auto x(a).$

I Know Some of You Have Been Wondering

▶ Consider:

```
auto x = type{value};
```

- ▶ Q: Does this “=” create a temporary object plus a move/copy?
 - ▶ Standard says “Yes, but”: The compiler may elide the temporary.
 - ▶ In practice, compilers do (and in the future routinely will) elide this temporary+move. However, the type must still be movable (which includes copyable as a fallback).

(The) Case Where You Can't Use “*auto Style*”

- ▶ **Case:** (1) Explicit “*type{}*” + (2) non-(cheaply-)moveable type.

```
auto lock = lock_guard<mutex>{ m }; // error, not movable  
auto ai = atomic<int>{};           // error, not movable  
auto a = array<int,50>{};          // compiles, but needlessly expensive
```

- ▶ **Non-cases:** Naked init list, proxy type, multi-word name.

```
auto x = { 1 }; // initializer_list  
auto x = 1;     // int  
auto a = matrix{...}, b = matrix{...}; // some lazily evaluated type  
auto ab = a * b; // capture proxy (efficient by default)  
auto c = matrix{ a * b }; // resolve computation  
auto x = (long long){ 42 }; // use int64_t{42} or 42LL  
auto y = class X{1,2,3}; // use X{1,2,3};
```

Cases Where You Can't ... Are Few

- ▶ A recent time I resisted using *auto*, I was wrong.

- ▶ It came up when changing this legacy code:

```
base* pb = new derived();
```

to this modern code, where I and others kept not noticing the different types:

```
unique_ptr<base> pb = make_unique<derived>();  
// too subtle: people keep not seeing it
```

and now I actually do prefer the consistent and nearly-as-terse spelling:

```
auto pb = unique_ptr<base>{ make_unique<derived>() };  
// explicit and clear: hard to miss it
```

which makes what's going on nice and explicit – the conversion is more obvious because we're explicitly asking for it.

Prefer declaring local variables using *auto*, whether the type should (1) track or (2) stick.

1. Deduced and exact, when you want tracking: `auto x = init`;
2. With explicit type name, when you want to commit: `auto x = Type { init }`;


Note: Guarantees zero implicit conversions/temporaries, zero narrowing conversions, and zero uninitialized variables!

Consider having some functions in headers (e.g., templates, inlines), return **auto** (only):
One-liners, and wrappers that should track type

They're in headers anyway. (Insert *de rigueur* modules note here.)

C++14 makes it convenient to not to not repeat yourself.

Remember: *auto* only \Rightarrow exact type, no conversions;
explicit return type \Rightarrow stable type, committed.



Remember, it's
really simple

To make type track, deduce:

```
auto var = init;  
auto f() { ... }
```

To make type stick, commit:

```
auto var = type{ init };  
auto f() -> type;  
or  
type var{ init };  
type f();
```

Use return-by-value way more often.

BUT: Don't overuse pass-by-value.

*Complete "how to pass params" details follow,
but the summary fits on a slide...
... one slide for "default," one slide for "optimal"*

Observation

"New features get overused." – B. Stroustrup

or

"It's about the lvalues, after all!" – S. Meyers

Just as exception safety isn't all about writing try and catch,
using move semantics isn't all about writing move and &&

Up Front: Acknowledgments & Hat Tips

- ▶ The following is the result of recent discussions with many people, including but not limited to the following:
 - ▶ Gabriel Dos Reis
 - ▶ **Matthew Fiovarante (&& param \equiv move from)**
 - ▶ **Howard Hinnant (distinguish copy ctor/op= costs vs. move)**
 - ▶ **Stephan T. Lavavej (low cost of value return even in C++98)**
 - ▶ **Scott Meyers (reduce #objects, be aware of costs)**
 - ▶ Eric Niebler
 - ▶ Sean Parent
 - ▶ **Bjarne Stroustrup (practicality, judgment, design sense)**
 - ▶ VC++ MVP discussion list
 - ▶ & many more



C++98: Reasonable Default Advice

	Cheap to copy (e.g., int)	Moderate cost to copy (e.g., string, BigPOD) or Don't know (e.g., unfamiliar type, template)	Expensive to copy (e.g., vector, BigPOD[])
Out	X f()		f(X&)*
In/Out	f(X&)		
In	f(const X&)		
In & retain copy	f(X)		

"Cheap" \approx a handful of hot int copies

"Moderate cost" \approx memcpy hot/contiguous \sim 1KB and no allocation

** or return X* at the cost of a dynamic allocation*

Modern C++: Reasonable Default Advice

	Cheap or impossible to copy (e.g., int, <code>unique_ptr</code>)	Cheap to move (e.g., <code>vector<T></code> , string) or Moderate cost to move (e.g., <code>array<vector></code> , BigPOD) or Don't know (e.g., unfamiliar type, template)	Expensive to move (e.g., <code>BigPOD[]</code> , <code>array<BigPOD></code>)
Out	X f()		f(X&)*
In/Out	f(X&)		
In	f(X)	f(const X&)	
In & retain "copy"			

Summary of what's new in C++1x:

✓ Defaults work better

* or return `unique_ptr<X>/make_shared<X>` at the cost of a dynamic allocation

Using the Advanced Knobs, Optimal

	Cheap or impossible to copy (e.g., int, <code>unique_ptr</code>)	Cheap to move (e.g., <code>vector<T></code> , string) or Moderate cost to move (e.g., <code>array<vector></code> , BigPOD) or Don't know (e.g., unfamiliar type, template)	Expensive to move (e.g., <code>BigPOD[]</code> , <code>array<BigPOD></code>)
Out	X f()		
In/Out	f(X&)		
In	f(X)	f(const X&)	
In & retain copy		f(const X&) + f(X&&) & move	
In & move from		f(X&&)	

+1 consistency: same optimization guidance as overloaded copy+move construction and assignment

Summary of what's new in C++1x:

✓ Defaults work better

✓ + More optimization opportunities

Using the Advanced Knobs, Optimal

	Cheap or impossible to copy (e.g., int, unique_ptr)	Cheap to move (e.g., vector<T>, string) or Moderate cost to move (e.g., array<vector>, BigPOD) or Don't know (e.g., unfamiliar type, template)	Expensive to move (e.g., BigPOD[], array<BigPOD>)
Out		X f()	f(X&)*
In/Out		f(X&)	
In		f(const X&)	
In & retain copy	f(X)	f(const X&) + f(X&&) & move**	
In & move from		f(X&&)**	

Summary of what's new in C++11x:

✓ Defaults work better

✓ + More optimization opportunities

When do I write rvalue &&? Only to optimize rvalues

Just as exception safety isn't all about writing try and catch, using move semantics isn't all about writing move and &&

Modern C++: A Narrowly Useful Option

	Cheap or impossible to copy (e.g., int, unique_ptr)	Cheap to move (e.g., vector<T>, string) or Moderate cost to move (e.g., array<vector>, BigPOD) or Don't know (e.g., unfamiliar type, template)	Expensive to move (e.g., BigPOD[], array<BigPOD>)
Out	X f()		f(X&)
In/Out	f(X&)		
In	f(const X&)		
In & retain copy	f(X)	f(X) & move ? *	
In & move from			

* GOOD: this can be faster than C++98 – can move from rvalues;

BUT: also can be much slower than C++98 – always incurs a full copy, prevents reusing buffers/state (e.g., for vectors & long strings, incurs memory allocation 100% of the time)

BUT: also problematic for noexcept

Journeyman Example: set_name

- ▶ Consider:

```
class employee {
    std::string name_;
public:
    void set_name( /*... ?? ...*/ ); // change name_ to new value
};
```

- ▶ Q: What should we tell people to write here?
 - ▶ Hint: There has been a lot of **overthinking** going on about this. (I include myself.)

Option #1: Default (same as C++98)

- ▶ Default: **const string&**

```
class employee {  
    std::string name_;  
public:  
    void set_name( const std::string& name ) { name_ = name; }  
};
```

- ▶ Always 1 copy assignment – *but usually* <<50% will alloc
 - ▶ If small (SSO), ~5 int copies, no mem alloc – often dominant
 - ▶ If large, still performs mem alloc <50% of the time

Option #2: Optimized (new for C++11)

- ▶ If optimization justified: Add overload for **string&&** + move

```
class employee {  
    std::string name_;  
public:  
    void set_name( const std::string& name ) { name_ = name; }  
    void set_name( std::string&& name ) noexcept  
        { name_ = std::move(name); }  
};
```

- ▶ Optimized to steal from rvalues:
 - ▶ Pass a named object: 1 copy assignment (<<50% alloc), as before
 - ▶ Pass a temporary: 1 move assignment (~5 ints, no alloc → **noexcept**)
 - ▶ Note: Combinatorial if multiple “in + retain copy” parameters.

Option #3: Pass by Value?

- ▶ Another new option in C++11: **string** + move

```
class employee {
    std::string name_;
public:
    void set_name( std::string name ) noexcept
                    { name_ = std::move(name); }
};
```

- ▶ Optimized to steal from rvalues, **without overloading**:
 - ▶ Pass named object: 1 copy *construction (100% alloc if long)* + move op=
 - ▶ Pass a temporary: 1 move assignment (~5 ints, no alloc → **noexcept-ish**)
 - ▶ This “noexcept” is... *problematic*

Option #4: Perfect Forwarding Idiom

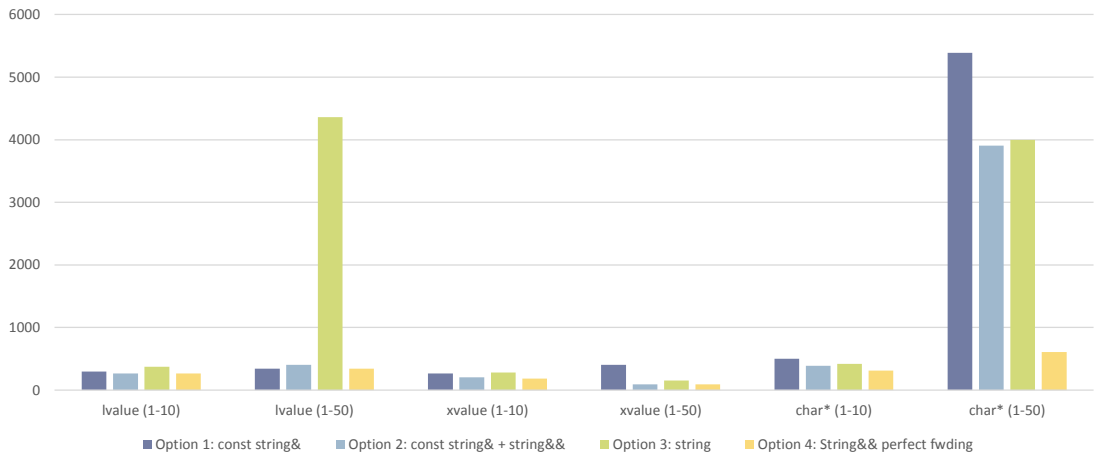
- ▶ Still another new option in C++11: **Templated T&&** “perfect forwarding”

```
class employee {
    std::string name_;
public:
    template<class String, class = std::enable_if_t<!std::is_same<std::decay_t<String>,
                                                    std::string::value>>
    void set_name( String&& name )
                    noexcept(std::is_nothrow_assignable<std::string&, String>::value)
                    { name_ = std::forward<String>(name); }
};
```

- ▶ Optimized to steal from rvalues (**and more**), **sort of without overloading**:
 - ▶ Pass a named object: 1 copy assignment (<<50% alloc), as before
 - ▶ Pass a temporary: 1 move assignment (~5 ints, no alloc → **noexcept**)
 - ▶ “Unteachable!” **Generates many funcs. Must be in a header. Can’t be virtual.**


Visual C++

VC++ 2013 x64 Release




Clang libc++

Clang/libc++ Release




gcc libstdc++

G++/libstdc++ x64 Release



gcc libstdc++ (coming soon)

G++/libstdc++ vstring x64 Release



vector & large string SMFs: Rough Costs

	Constructor	operator=
Default	\$\$	
Move		\$
Copy	\$\$\$\$	\$\$\$

(More) Geek Heroes

- ▶ Howard Hinnant: **“Don’t blindly assume that the cost of construction is the same as assignment.”**
 - ▶ For strings and vectors, “Capacity plays a large role in their performance. Copy construction always allocates (except for short). Copy assignment (except for short) allocates/deallocates 50% of the time with random capacities on the lhs and rhs. To keep an eye on performance, one must count allocations and deallocations.”
- ▶ William of Occam: **‘Do not multiply entities needlessly.’**
 - ▶ Attributed. Talking about hypotheses; applies to ‘entities.’
- ▶ Andrei Alexandrescu: **“No work is less work than some work.”**
- ▶ Scott Meyers: **‘It’s a bad habit to just create extra objects.’**
 - ▶ “Just create ‘em because they’re cheap to move from” is thoughtcrime.



This Talk

REPRISE

- ▶ This talk focuses on **defaults**, basic styles and idioms in modern C++.
 - ✗ ~~“Default” != “don’t think.”~~
 - ✓ “Default” == “don’t **overthink.**” Esp. don’t **optimize prematurely.**
- ▶ These reinforce (not compete with) the “fundamentals.”
 - ▶ **“Write for clarity and correctness first.”**
 - ▶ **“Avoid premature optimization.”** By default, prefer *clear* over *optimal*.
 - ▶ **“Avoid premature pessimization.”** Prefer *faster* when equally *clear*.

Option #3: Pass by Value?

- ▶ Another new option in C++11: **string** + move

```
class employee {  
    std::string name_;  
public:  
    void set_name( std::string name ) noexcept  
                    { name_ = std::move(name); }
```

An interesting attempt that temporarily drew in a number of experts!
But: at least “too cute” & probably just an antipattern... *except for one case...*

overloading:
100% alloc if long) + move op=
ints, no alloc → **noexcept-ish**)

Option #3: Pass by Value *for Constructors*

- ▶ There is one place where this is a good idea: **Constructors**.

```
class employee {  
    std::string name_;  
    std::string addr_;  
    std::string city_;  
public:  
    void employee( std::string name, std::string addr, std::string city )  
        : name_{std::move(name)}, addr_{std::move(addr)}, city_{std::move(city)} {}  
};
```

- ▶ Constructors are the primary case of **multiple** “in + retain copy” params, where overloading const&/&& is combinatorial.
- ▶ Constructors always construct, so no worries about **reusing** existing capacity.
- ▶ Note: Probably prefer not to write the misleading “noexcept” ...

Option #1: **Default** (same as C++98)

- ▶ Default: **const string&**

```
class employee {  
    std::string name_;  
public:  
    void set_name( const std::string& name ) { name_ = name; }  
};
```

- ▶ Always 1 copy assignment – *but usually* <<50% will alloc
 - ▶ If small (SSO), ~5 int copies, no mem alloc – often dominant
 - ▶ If large, still performs mem alloc <50% of the time

Option #2: **Optimized** (new for C++11)

- ▶ **If optimization justified:** Add overload for **string&&** + move

```
class employee {  
    std::string name_;  
public:  
    void set_name( const std::string& name ) { name_ = name; }  
    void set_name( std::string&& name ) noexcept  
        { name_ = std::move(name); }  
};
```

- ▶ Optimized to steal from rvalues:
 - ▶ Pass a named object: 1 copy assignment (<<50% alloc), as before
 - ▶ Pass a temporary: 1 move assignment (~5 ints, no alloc → **noexcept**)
 - ▶ Note: Combinatorial if multiple “in + retain copy” parameters.

What is a T&&? A forwarding reference


Quiz

```
void foo( X&& x );  
template<class Y>  
void bar( Y&& y );
```

- ▶ **Q: What are the types of the function parameters?
What arguments to they accept or reject?
What is the parameter for?**
- ▶ **A: Fundamentally different.**
 - ▶ *foo* takes **rvalue** reference to **non-const**.
foo accepts **only rvalue** X objects.
foo's parameter is to **capture temporaries** (and other rvalues).
 - ▶ *bar* takes *mumble* reference to **everything**: const, volatile, both, **and** neither.
bar accepts **all** Y objects.
bar's parameter is for **forwarding** its argument onward.

Forwarding References

- ▶ Scott Meyers pointed out that T&& is very different, and needs a name.
 - ▶ He coined "universal reference."
 - ▶ For his book whose final galleys are due, um, today.
- ▶ Here at CppCon, a few of us met and ultimately agreed that this does need a name. (Thanks, Scott.)
 - ▶ But we still disliked "universal." (Sorry, Scott.)
 - ▶ We think the right name is "**forwarding reference**."
 - ▶ The committee/community may disagree. Time will tell.
 - ▶ In the meantime, Scott will add a footnote and index entry for "forwarding reference," and switch to it in future printings if the community agrees. (Thanks, Scott!)



Uses and Abuses of &&

- ▶ Use && only for parameter/return types:
 - ▶ **myclass&& rvalue references to optimize rvalues**, usually overloading const& / && – note this covers the move SMFs!

```
void f( const string& );    // default way to express “in + retain a copy”  
void f( string&& );       // what to add to additionally optimize for rvalues
```
 - ▶ **T&& forwarding references to write forwarders**, which are neutral code between unknown callers and callees and want to preserve rvalueness/cv-ness.
 - ▶ Note this includes the new proposed *for(e:c)*, which is... *drum roll...* a neutral forwarder between a collection/range and the calling code.
 - ▶ Also includes generic lambda auto&& parameters... use for forwarders only.
- ▶ Don't use auto&& for local variables.
 - ▶ You should know whether your variable is const/volatile or not!
 - ▶ (Except rarely if you're just handing it off... in the body of a forwarder.)

Dessert Slide:

Use **tuple** for multiple return values.

Yes, C++11 has multiple return values! (Who knew?)

Sweet Realization: We're Already Doing It

- ▶ Given a `set<string>` `myset`, consider:

```
// C++98
```

```
pair<set<string>::iterator,bool> result = myset.insert( "Hello" );  
if (result.second) do_something_with( result.first );    // workaround
```

```
// C++11 – sweet backward compat
```

```
auto result = myset.insert( "Hello" );                // nicer syntax, and the  
if (result.second) do_something_with( result.first );  // workaround still works
```

```
// C++11 – sweet forward compat, can treat as multiple return values
```

```
tie( iter, success ) = myset.insert( "Hello" );       // normal return value  
if (success) do_something_with( iter );
```

It's hard to *be-remember you're* an expert

C++ developers (~3M)

REPRISE

libstdc++ developers (~30)
+
libc++ developers (~5-7)
+
Boost developers (~300?)
+
ISO WG21 attendees (~300?)

Back to Basics: Modern C++ Style

loops
pointers & references
smart pointers
variable declarations
parameter passing

Questions?