

# HOW UBISOFT MONTREAL DEVELOPS GAMES FOR MULTICORE

*Before & After C++11*

Jeff Preshing  
Technical Architect  
Ubisoft Montreal

UBISOFT

CAFÉ VIENNE

CAFÉ VIENNE



108.4

Ulysse-service

Super Ulysse

Depot

Super Ulysse

Depot

**Watch Dogs**

**Assassin's Creed Unity**

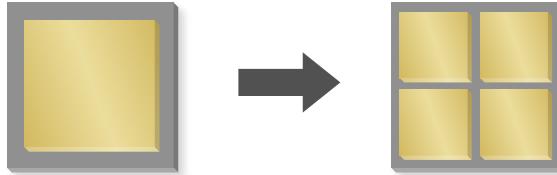
# MULTICORE

**Rainbow Six: Siege**

**Far Cry 4**

# MULTICORE

*several CPU cores on a single processor*



- Same instruction set
- Same address space
- Existing threads can run on any core

*Popular way to offer **power***



**Game  
Industry**

**C++  
Community**



*We all want to exploit multicore!*

# PART ONE

*Multicore Programming at Ubisoft*

# PART TWO

*The C++11 Atomic Library*

Part One

# Multicore Programming at Ubisoft

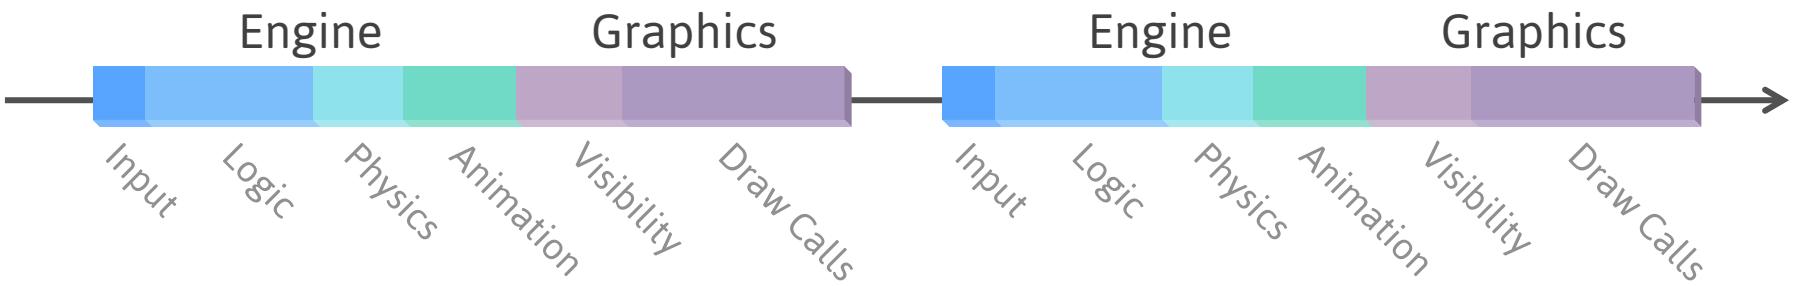
# GENERAL-PURPOSE HARDWARE THREADS

*available for game use*



# SINGLE-THREADED MAIN LOOP

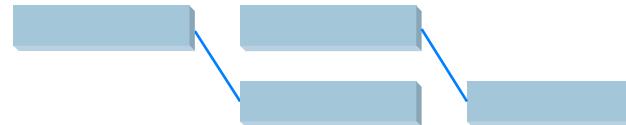
*in the early 2000s*



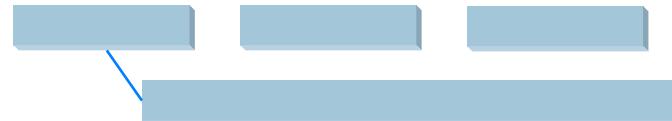
# THREE THREADING PATTERNS

*to exploit multicore*

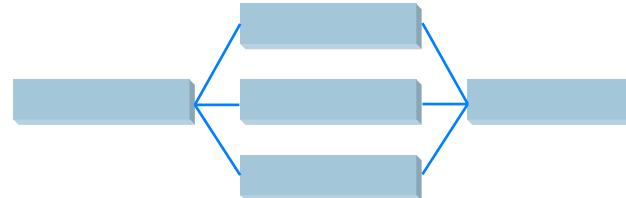
Pipelining Work



Dedicated Threads

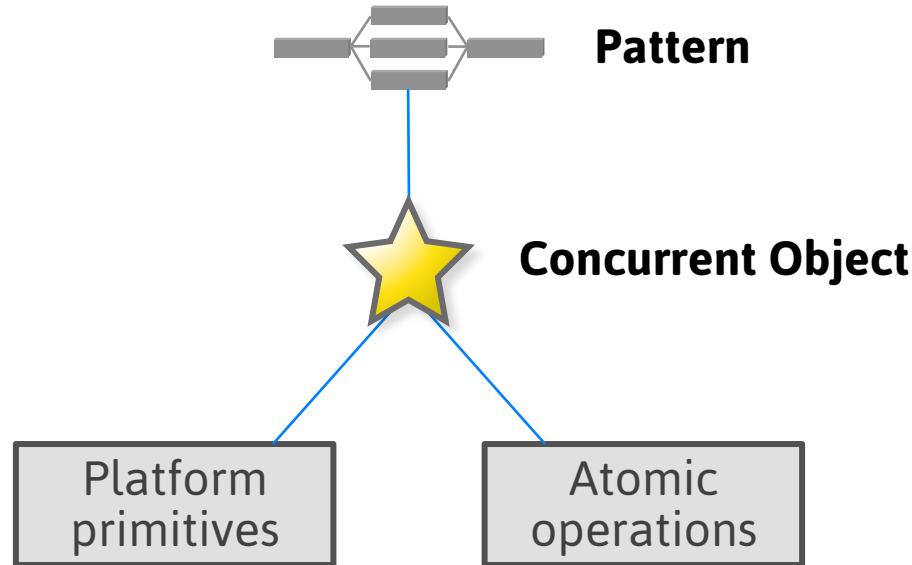


Task Schedulers



# CONCURRENT OBJECTS

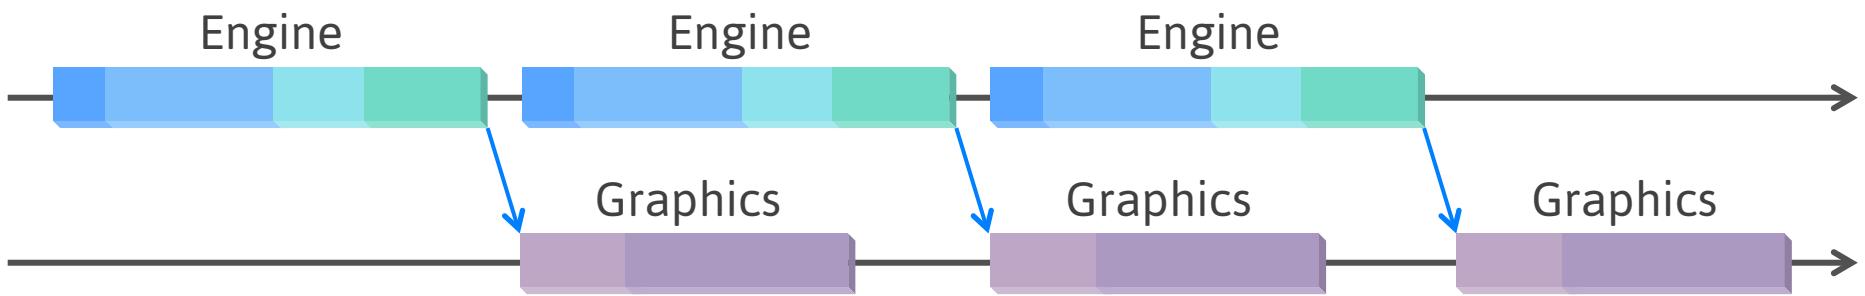
Multiple threads operate concurrently on the object with at least one thread modifying its state.



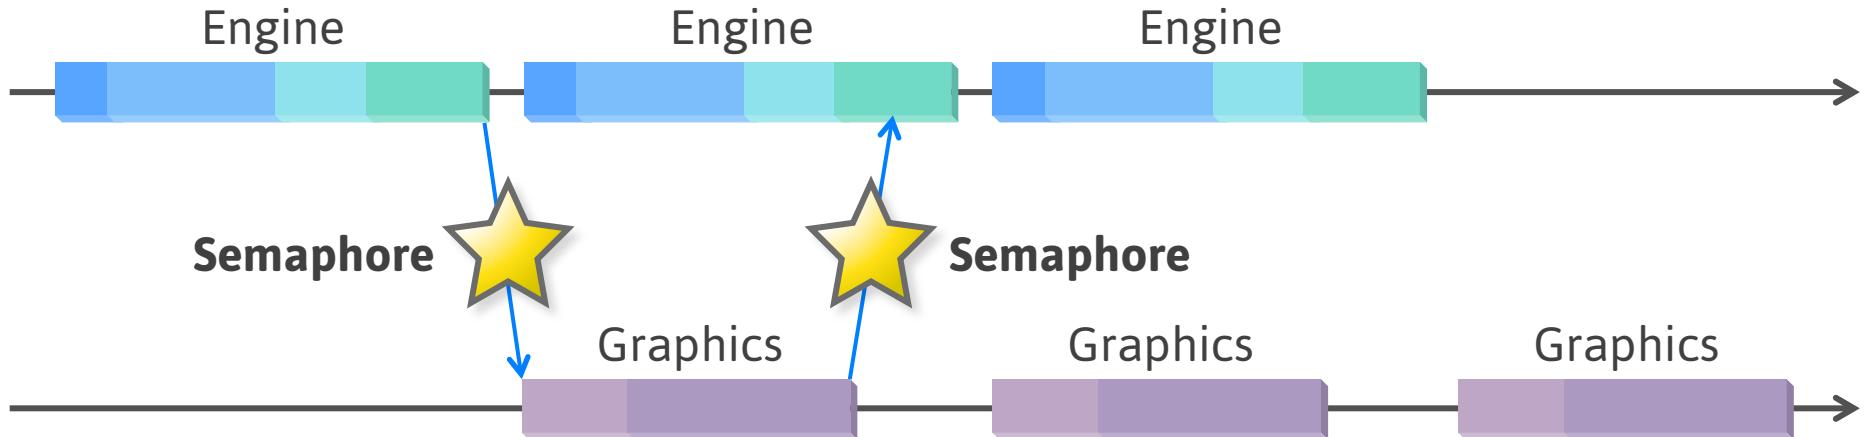
# Pipelining Work



# PIPELINED GRAPHICS

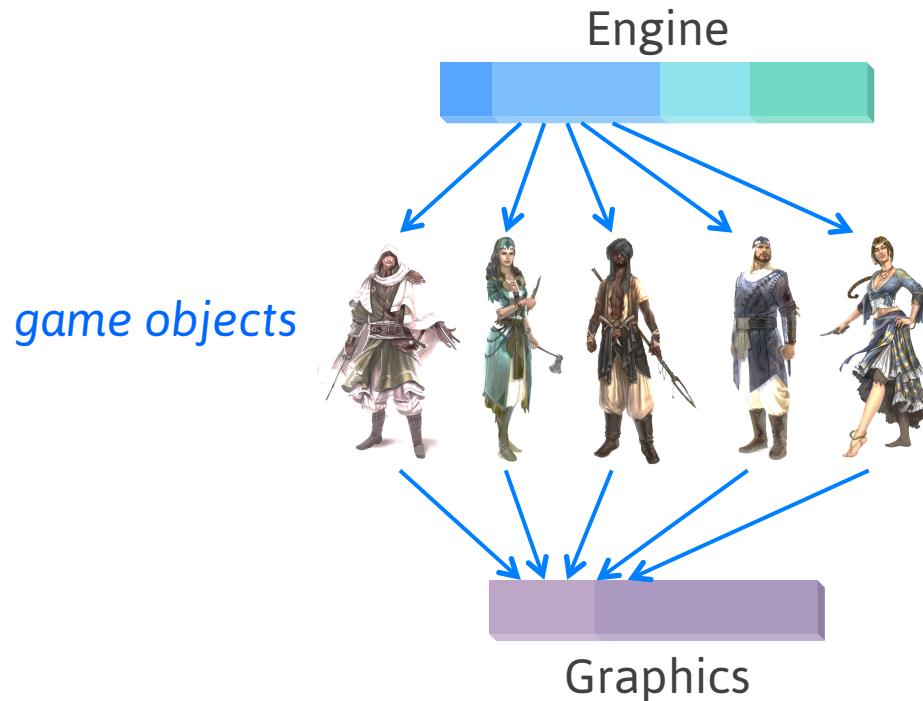


# PIPELINED GRAPHICS



# PIPELINED GRAPHICS

how to avoid concurrent modifications?



# DOUBLE-BUFFERED GRAPHICS STATE

*One approach*



```
struct Object
{
    ...
    Matrix xform[2];
    ...
};
```

# SEPARATE GRAPHIC OBJECTS

*Another approach*



```
struct Object
{
    ...
    Matrix xform; ——————>
    GraphicObject* gfxObject;
    ...
};
```



```
struct GraphicObject
{
    ...
    Matrix xform;
    ...
};
```

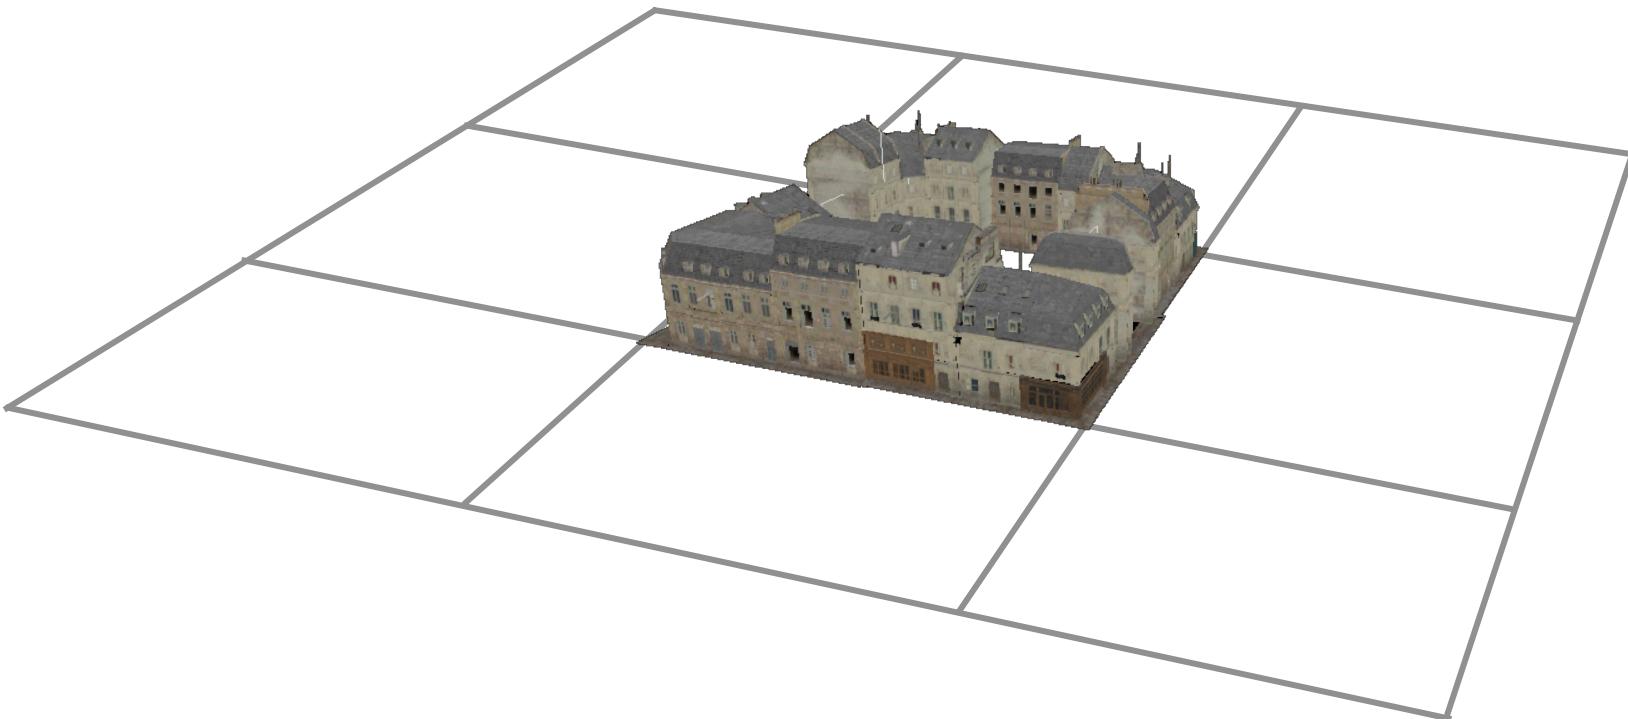
Copy at start of frame

# Dedicated Threads

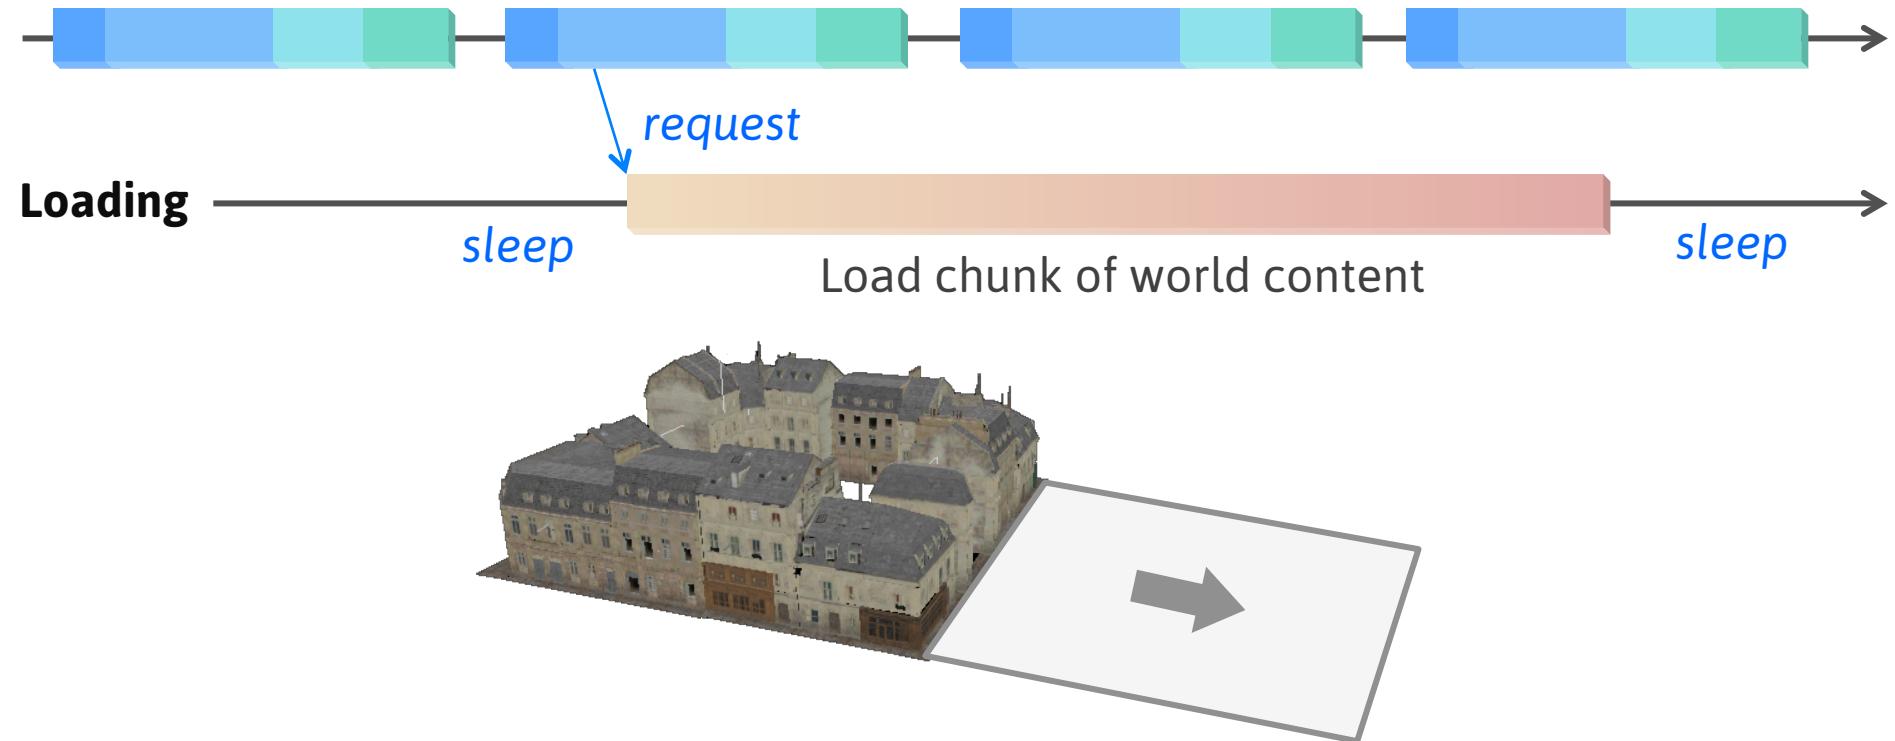


# CONTENT STREAMING

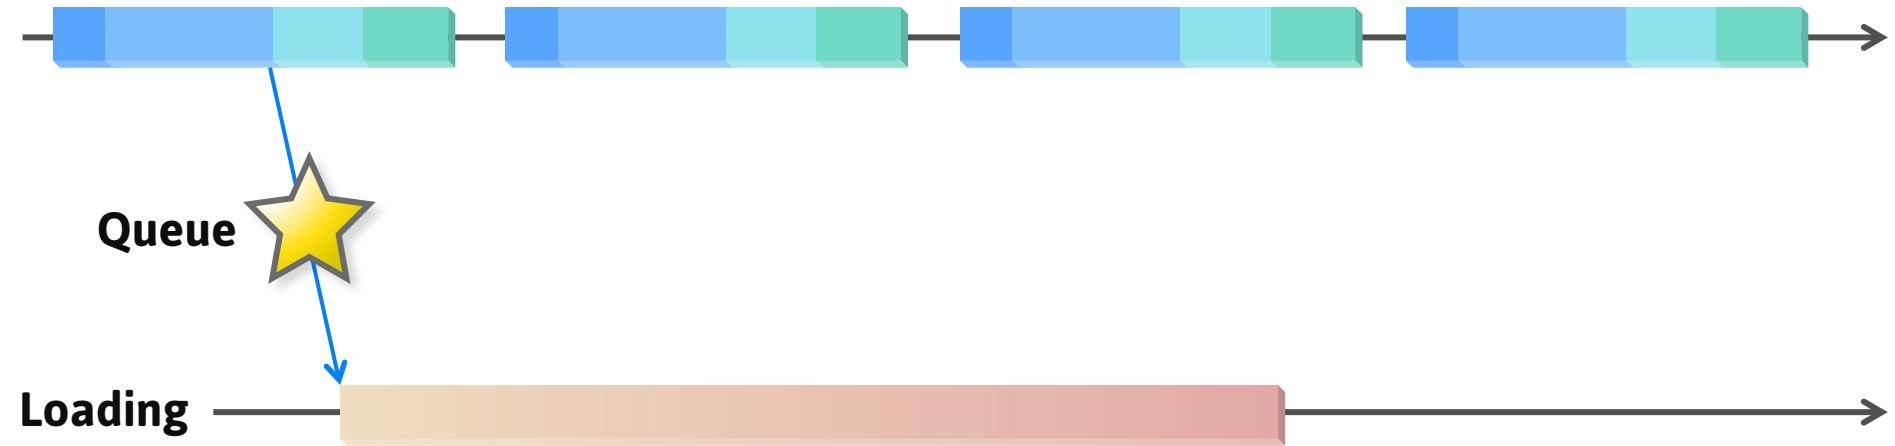
We don't load the entire game environment in memory at once.



# DEDICATED LOADING THREAD



# DEDICATED LOADING THREAD



# WAKING UP THE LOADING THREAD

```
ThreadSafeQueue<Request> requests;  
Event workAvailable;
```

## Engine Thread

```
requests.push(r);  
workAvailable.signal();
```

## Loading Thread

```
for (;;)  
{  
    workAvailable.waitAndReset();  
    while (r = requests.tryPop())  
    {  
        processLoadRequest(r);  
    }  
}
```

Event signaled → threads pass through

Event reset → threads wait

# IMPROVING ON THE QUEUE

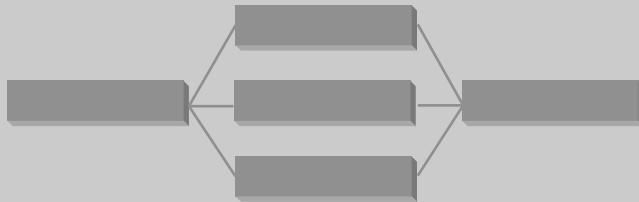
*Many design choices*

Custom



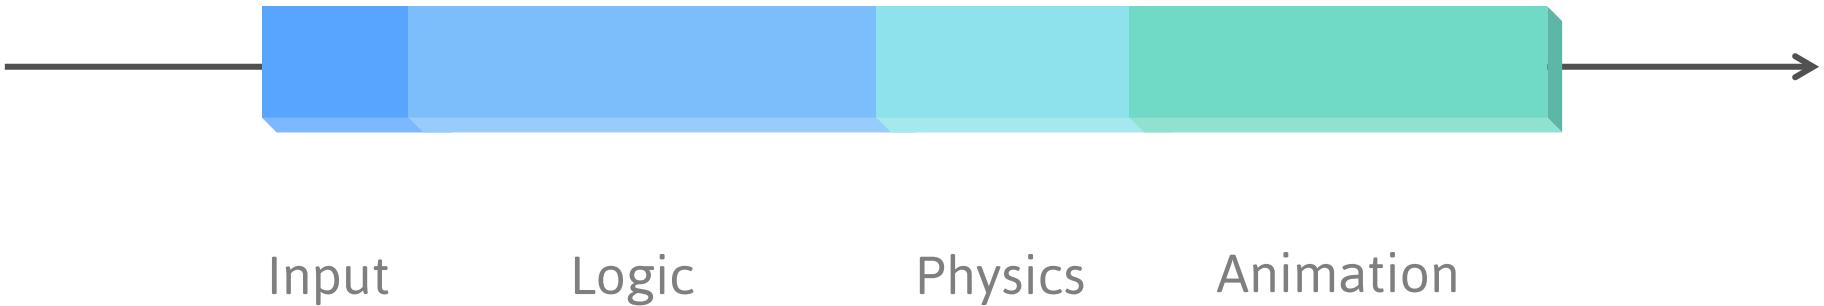
- Cancel requests
- Interrupt requests
- Re-prioritize requests

# Task Schedulers



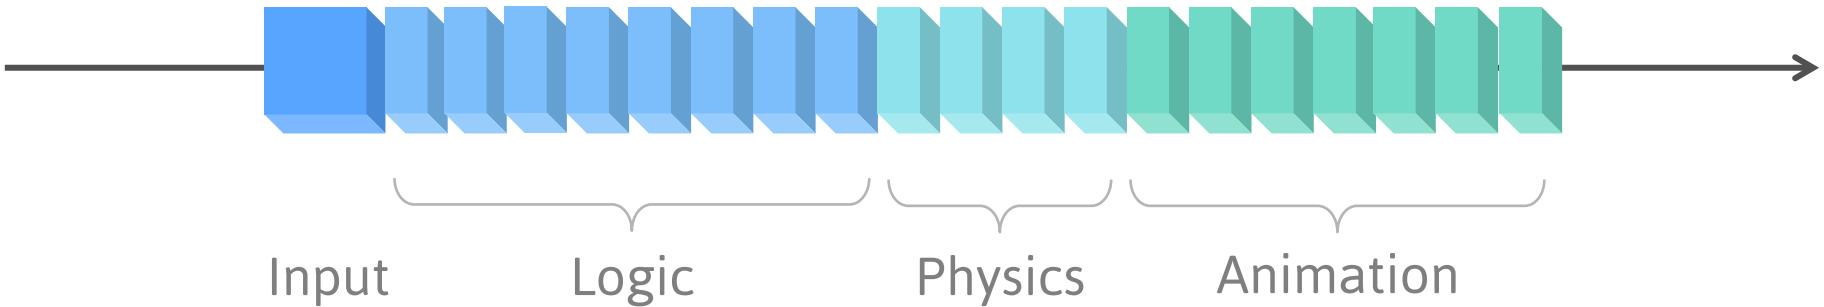
# FINE-GRAINED PARALLELISM

*Motivation for a task scheduler*



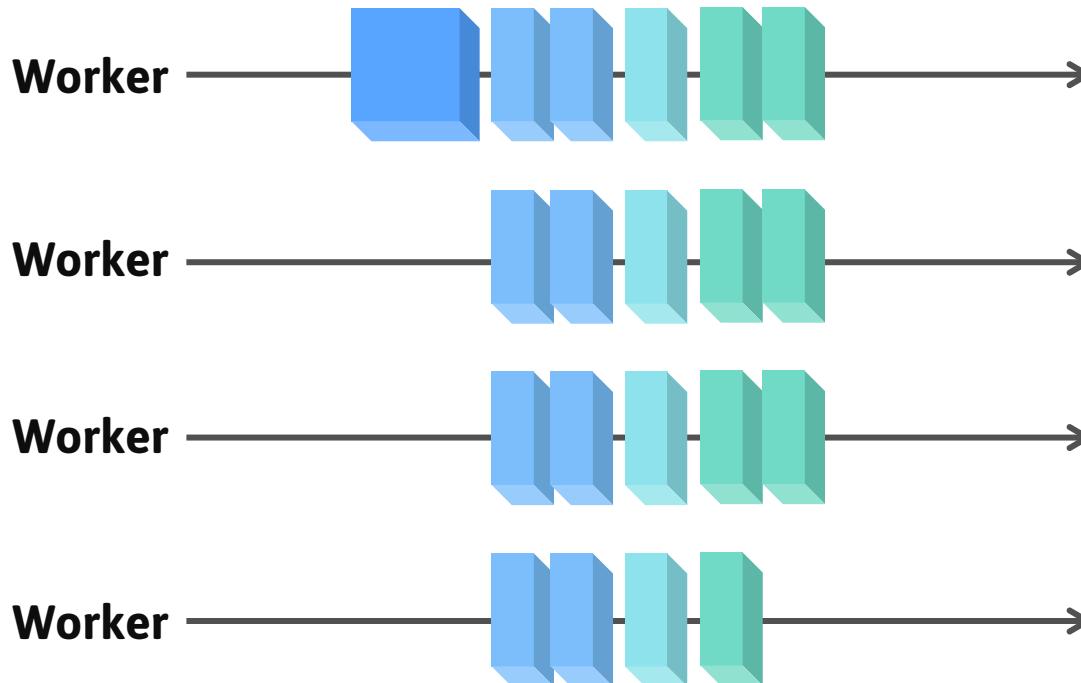
# FINE-GRAINED PARALLELISM

*Motivation for a task scheduler*

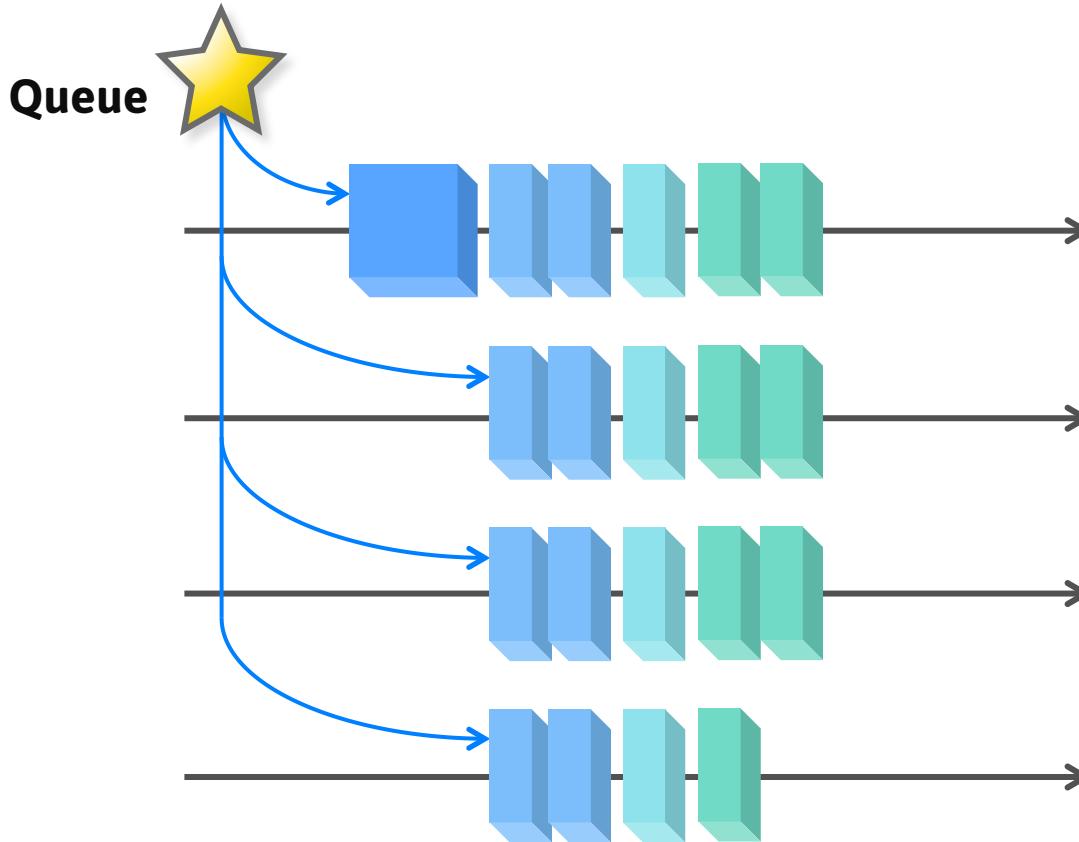


# FINE-GRAINED PARALLELISM

*Motivation for a task scheduler*



# SIMPLE TASK QUEUE



# WAKING UP THE WORKER THREADS

```
ThreadSafeQueue<Task> tasks;  
Event workAvailable[numThreads];
```

## Submitting Thread

```
tasks.push(t);  
for (int i = 0; i < numThreads; i++)  
    workAvailable.signal();
```

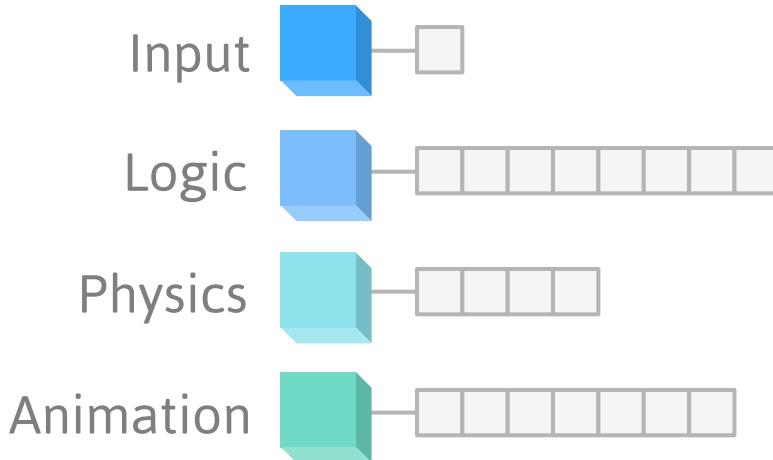
## Worker Thread

```
for (;;)  
{  
    workAvailable[thread].waitAndReset();  
    while (t = tasks.tryPop())  
    {  
        t->Run();  
    }  
}
```

One event for each worker thread

# TASK GROUPS

*Grouping work units together into larger tasks*



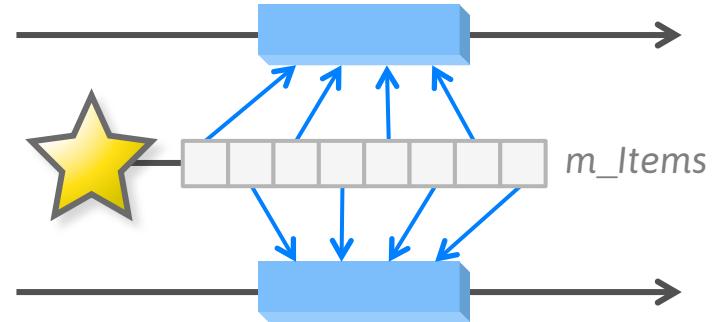
```
class TaskGroup
{
private:
    Array<Item*> m_Items;
    ...
};
```

Each TaskGroup keeps an array of items to update in parallel.

# TASK GROUPS

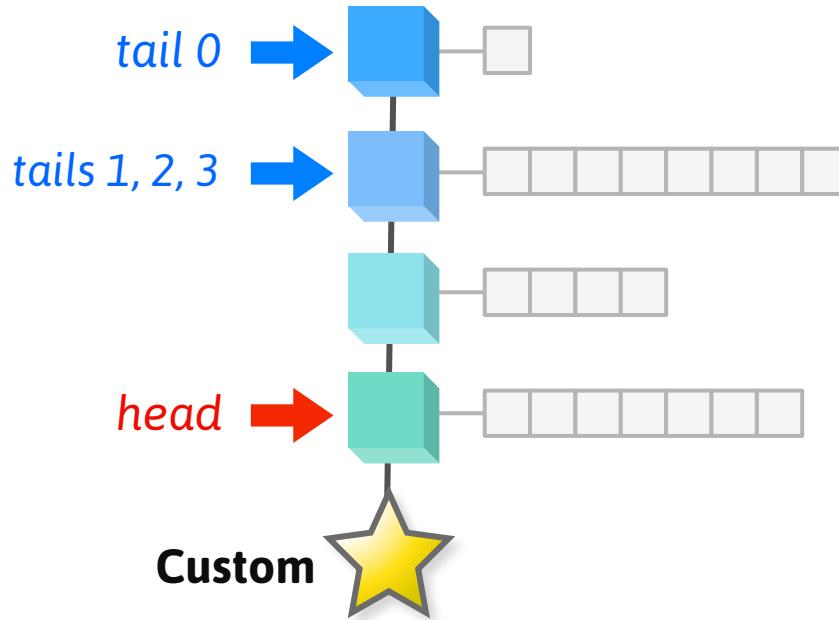
```
class TaskGroup
{
private:
    vector<Item*> m_Items;
    volatile int m_Index;

public:
    void Run()
    {
        for (;;)
        {
            int index = AtomicIncrement(m_Index);
            if (index >= m_Items.size())
                break;
            m_Items[index]->Run();
        }
    }
};
```



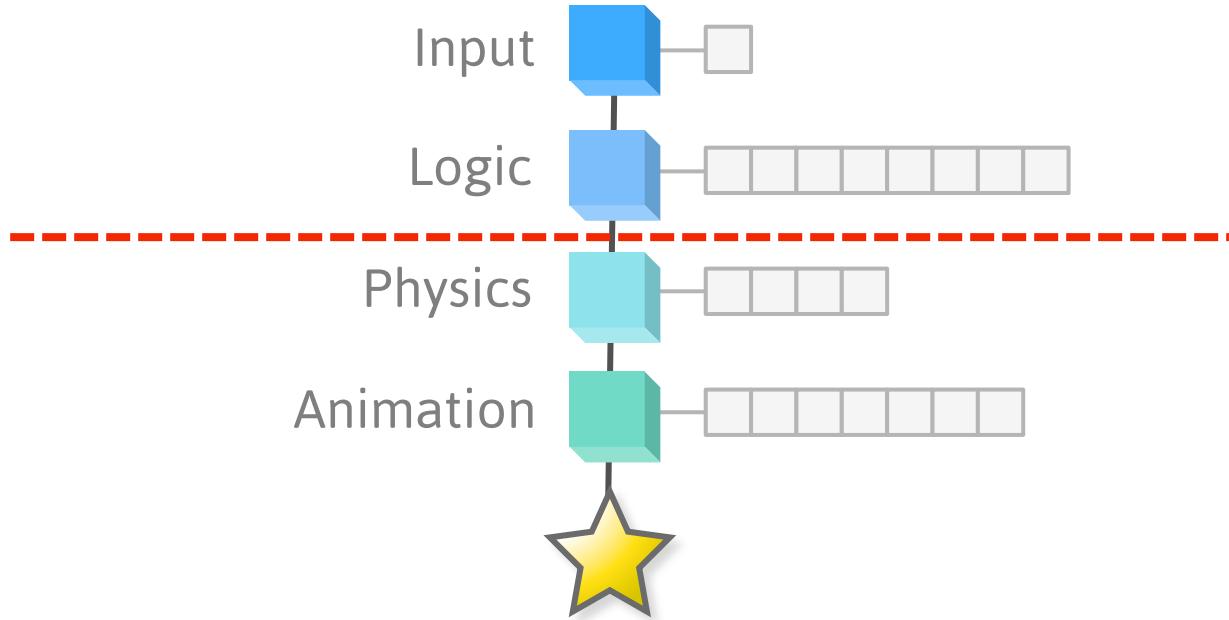
Multiple threads work on the same TaskGroup.

# NOT A SIMPLE QUEUE ANYMORE



Could be a queue with separate tails for each worker.

# MANAGING DEPENDENCIES

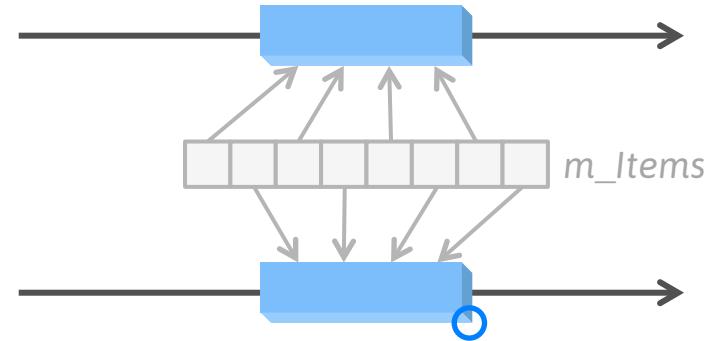


**No** physics tasks before ***all*** logic tasks.

# MANAGING DEPENDENCIES

```
class TaskGroup
{
private:
    vector<Item*> m_Items;
    volatile int m_Index;
    volatile int m_RemainingCount;
    ...

public:
    void Run() {
        int count = 0;
        for (;;) {
            int index = AtomicIncrement(m_Index);
            if (index >= m_Items.size())
                break;
            m_Items[index]->Run();
            count++;
        }
        if (count > 0 && AtomicSubtract(m_Index, count) == 0)
            AddDependencies();
    }
};
```



The thread that finishes the last item schedules the next TaskGroup.

# IMPROVING ON THE TASK SCHEDULER

*Many design choices*

Custom



- Centralized / per-thread task list
- Priorities
- Affinities
- Batching
- Profiler integration
- “Pin” threads to cores

# Atomic Operations

# GAME ATOMICS

*Typical portable library*

<b>Declaration</b>	volatile int A;
<b>Load/Store</b>	A = 1; int a = A;
<b>Ordering</b>	LIGHTWEIGHT_FENCE(); FULL_FENCE();
<b>Read-Modify-Write</b>	AtomicIncrement(A); AtomicCompareExchange(A, ..., ...); ...

# FENCE MACROS

*What's the difference?*

**Used more often**

**LIGHTWEIGHT\_FENCE();**

Orders loads from memory

Orders stores to memory

Does the job of:

```
atomic_thread_fence(memory_order_acquire);  
atomic_thread_fence(memory_order_release);  
atomic_thread_fence(memory_order_acq_rel);
```

**FULL\_FENCE();**

*... all that, plus:*

Commits stores before next load

Does the job of:

```
atomic_thread_fence(memory_order_seq_cst);
```

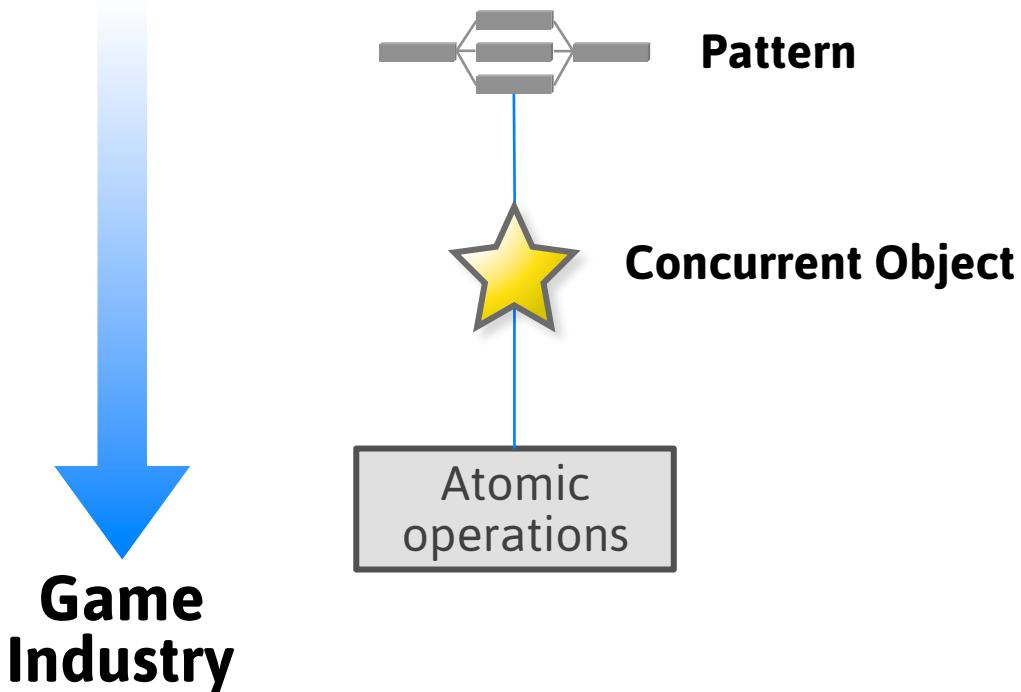
# HOW THEY'RE IMPLEMENTED

*on processors we care about*

		x86/64	PowerPC	ARMv7
<b>Declaration</b>	volatile int A;			
<b>Load/Store</b>	A = 1; int a = A;	mov %, % mov %, %	ld %, % st %, %	ldr %, % str %, %
<b>Ordering</b>	LIGHTWEIGHT_FENCE(); FULL_FENCE();	(compiler barrier) mfence	lwsync hwsync	dmb dmb
<b>Read-Modify-Write</b>	AtomicIncrement(A); AtomicCompareExchange(A, ..., ...); ...	lock inc lock cmpxchg	lwarx ... stwcx	ldrex ... strex

# ATOMIC OPERATIONS

*How we ended up using them*

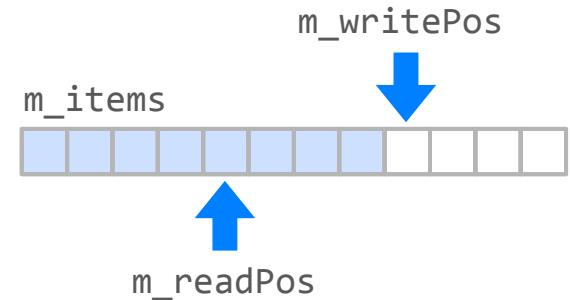


# EXAMPLE

## Capped wait-free queue

```
template <class T, int size>
class CappedSPSCQueue
{
private:
    T m_items[size];
    volatile int m_writePos;
    int m_readPos;

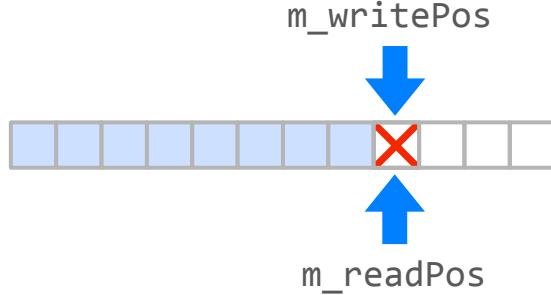
public:
    CappedSPSCQueue() : m_writePos(0), m_readPos(0) {}
    bool tryPush(const T& item) { ... }
    bool tryPop(T& item) { ... }
};
```



Single producer, single consumer

# EXAMPLE

## Capped wait-free queue



```
bool tryPush(const T& item)
{
    int w = m_writePos;
    if (w >= size)
        return false;
    m_items[w] = item;
    m_writePos = w + 1;
    return true;
}
```

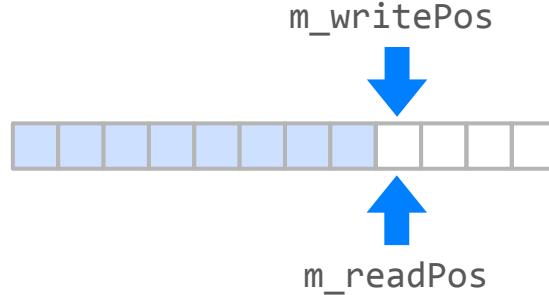
**BROKEN**

```
bool tryPop(T& item)
{
    int w = m_writePos;
    if (m_readPos >= w)
        return false;
    item = m_items[m_readPos];
    m_readPos++;
    return true;
}
```

A red dashed oval encloses the code for `tryPush`. A blue curved arrow labeled "reorder" points from the `tryPush` code to the `tryPop` code, indicating a race condition or bug where the `tryPop` operation might see the item pushed by the `tryPush` operation before it has been written.

# EXAMPLE

## Capped wait-free queue



```
bool tryPush(const T& item)
{
    int w = m_writePos;
    if (w >= size)
        return false;
    m_items[w] = item;
    LIGHTWEIGHT_FENCE();
    m_writePos = w + 1;
    return true;
}
```

```
bool tryPop(T& item)
{
    int w = m_writePos;
    if (m_readPos >= w)
        return false;
    LIGHTWEIGHT_FENCE();
    item = m_items[m_readPos];
    m_readPos++;
    return true;
}
```

# RECAP:

## Multicore programming at Ubisoft

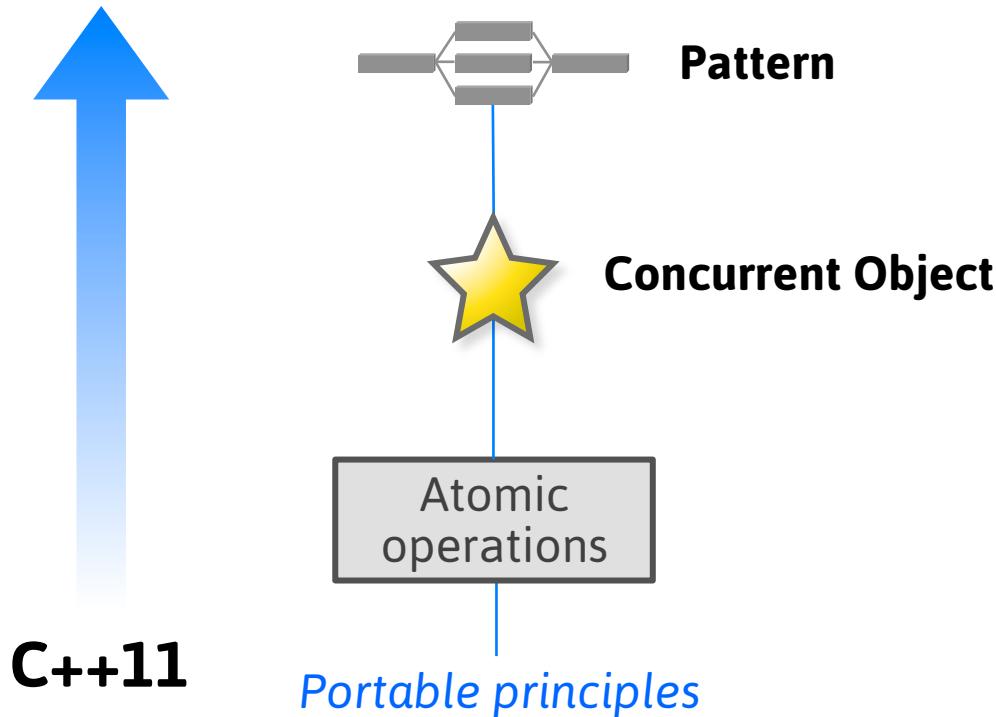
- Three threading patterns
- Lots of custom concurrent objects
- Atomic operations for high contention objects
- We learned by doing

Part Two

# The C++11 Atomic Library

# ATOMIC OPERATIONS

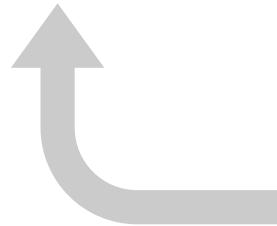
*in C++11*



# C++11 FORBIDS DATA RACES

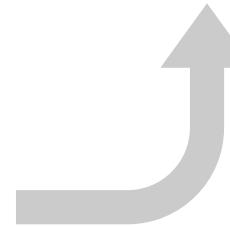
If multiple threads access the same variable concurrently, and at least one thread modifies it, all threads must use C++11 atomic operations.

Thread 1



```
int X;
```

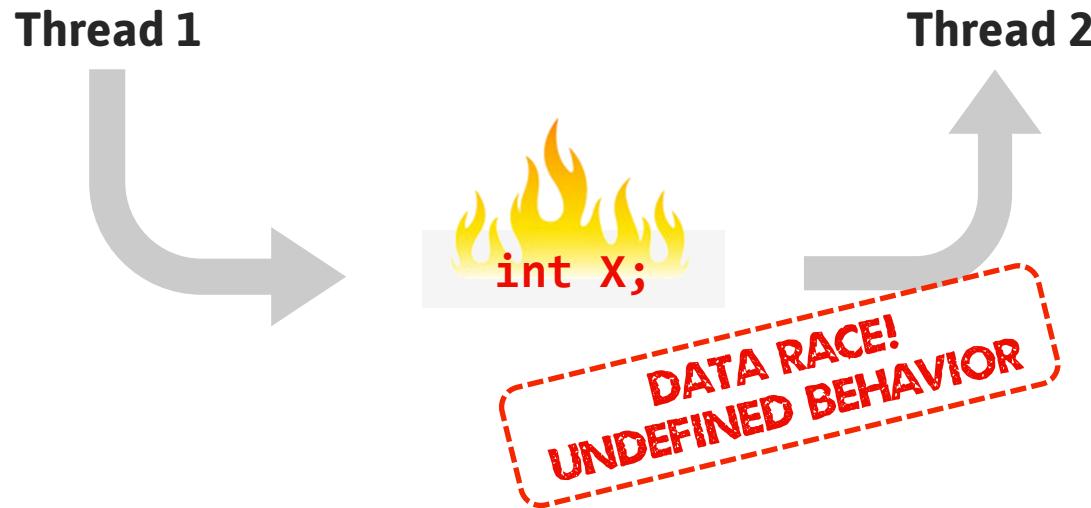
Thread 2



OK!

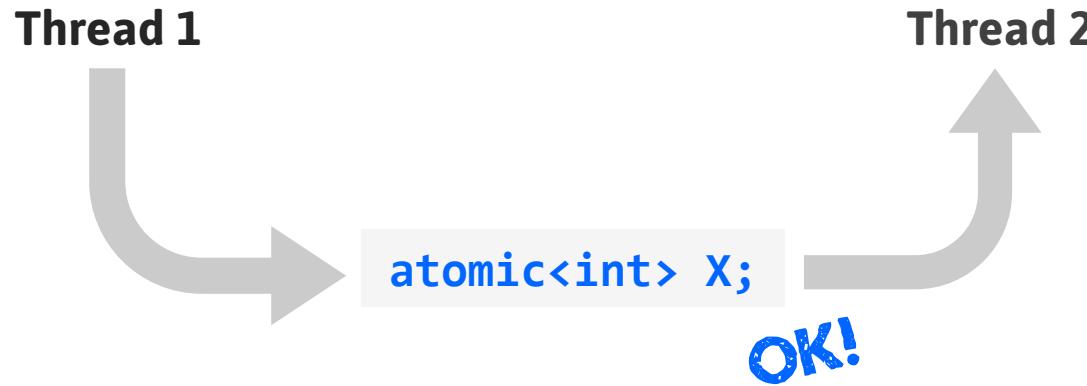
# C++11 FORBIDS DATA RACES

If multiple threads access the same variable concurrently, and at least one thread modifies it, all threads must use C++11 atomic operations.



# C++11 FORBIDS DATA RACES

If multiple threads access the same variable concurrently, and at least one thread modifies it, all threads must use C++11 atomic operations.



- That's how you know when you must use `atomic<>`.

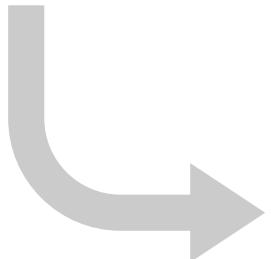
# C++11 FORBIDS DATA RACES

*One reason they're bad*

**Thread 1**

```
x = 0x80004;
```

*If machine can  
only write 16 bits...*



**Thread 2**

```
c = x;
```

*...we get a "torn write".  
0x80000*

`int X;`  
*...and this is 32-bit...*

# C++11 FORBIDS DATA RACES

If multiple threads access the same variable concurrently, and at least one thread modifies it, all threads must use C++11 atomic operations.



We break this rule all the time.  
We know that `int` is atomic.

# IT'S ACTUALLY TWO ATOMIC LIBRARIES

*Masquerading under one API*

## Sequentially Consistent Atomics

- Similar to Java volatiles
- Used in literature/books

## Low-Level Atomics

- Similar to C/C++ volatiles
- Much like game atomics

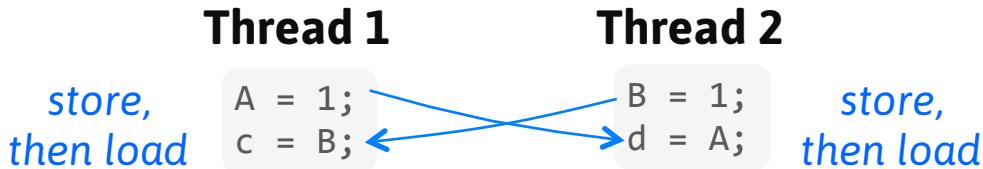
All about **interleaving** statements



# SEQUENTIALLY CONSISTENT ATOMICS

## Example #1

```
atomic<int> A(0);  
atomic<int> B(0);
```



## Possible Interleavings:

A = 1;  
c = B;  
B = 1;  
d = A;

A = 1;  
B = 1;  
c = B;  
d = A;

B = 1;  
d = A;  
A = 1;  
c = B;

c	d
0	0
0	1
1	0
1	1

Impossible!

# LOW-LEVEL ATOMICS

## Example #1

```
atomic<int> A(0);  
atomic<int> B(0);
```

### Thread 1

```
A.store(1, memory_order_relaxed);  
c = B.load(memory_order_relaxed);
```

### Thread 2

```
B.store(1, memory_order_relaxed);  
d = A.load(memory_order_relaxed);
```

*Doing the same thing*

You can prevent it with “full memory fences”:

```
atomic_thread_fence(memory_order_seq_cst);
```

c	d
0	0
0	1
1	0
1	1

*Possible!*

# SEQUENTIALLY CONSISTENT ATOMICS

How to write them

```
atomic<int> A;
```

```
A.store(1, memory_order_seq_cst); ← All other constraints are low-level  
c = A.load(memory_order_seq_cst);
```

...is the same as:

```
A.store(1);  
c = A.load();
```

Default argument

...and the same as:

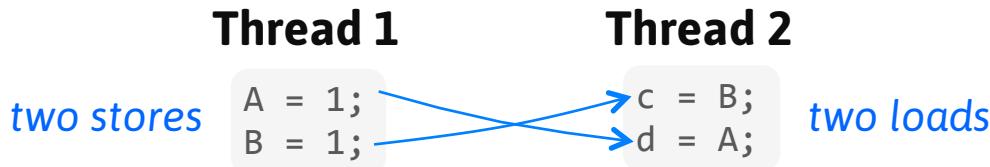
```
A = 1;  
c = A;
```

Operator overloading

# SEQUENTIALLY CONSISTENT ATOMICS

## Example #2

```
atomic<int> A(0);  
atomic<int> B(0);
```



## Possible Interleavings:

```
A = 1;  
B = 1;  
c = B;  
d = A;
```

```
A = 1;  
c = B;  
B = 1;  
d = A;
```

```
c = B;  
d = A;  
A = 1;  
B = 1;
```

c	d
0	0
0	1
1	0
1	1

Impossible!

# LOW-LEVEL ATOMICS

## Example #2

```
atomic<int> A(0);  
atomic<int> B(0);
```

### Thread 1

```
A.store(1, memory_order_relaxed);  
B.store(1, memory_order_relaxed);
```

### Thread 2

```
c = B.load(memory_order_relaxed);  
d = A.load(memory_order_relaxed);
```

*Doing the same thing*

This is the bug from Section One!

You can fix it with “lightweight fences”:

```
atomic_thread_fence(memory_order_acquire);  
atomic_thread_fence(memory_order_release);
```

c	d
0	0
0	1
1	0
1	1

Possible!

# VISUALIZING LOW-LEVEL ATOMICS

Imagine each thread having its own private copy of memory.

Thread 1

A	1
B	0

```
A.store(1, memory_order_relaxed);  
c = B.load(memory_order_relaxed);
```

Thread 2

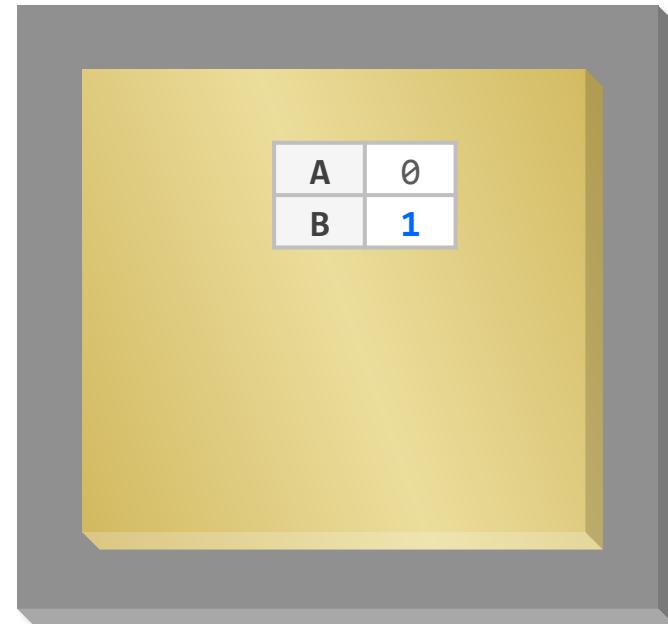
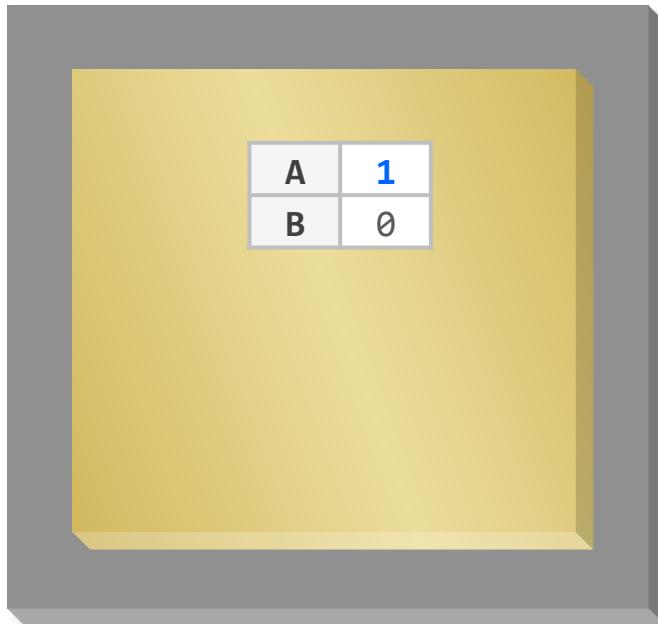
A	0
B	1

```
B.store(1, memory_order_relaxed);  
d = A.load(memory_order_relaxed);
```

Now c = 0, d = 0 is trivial.

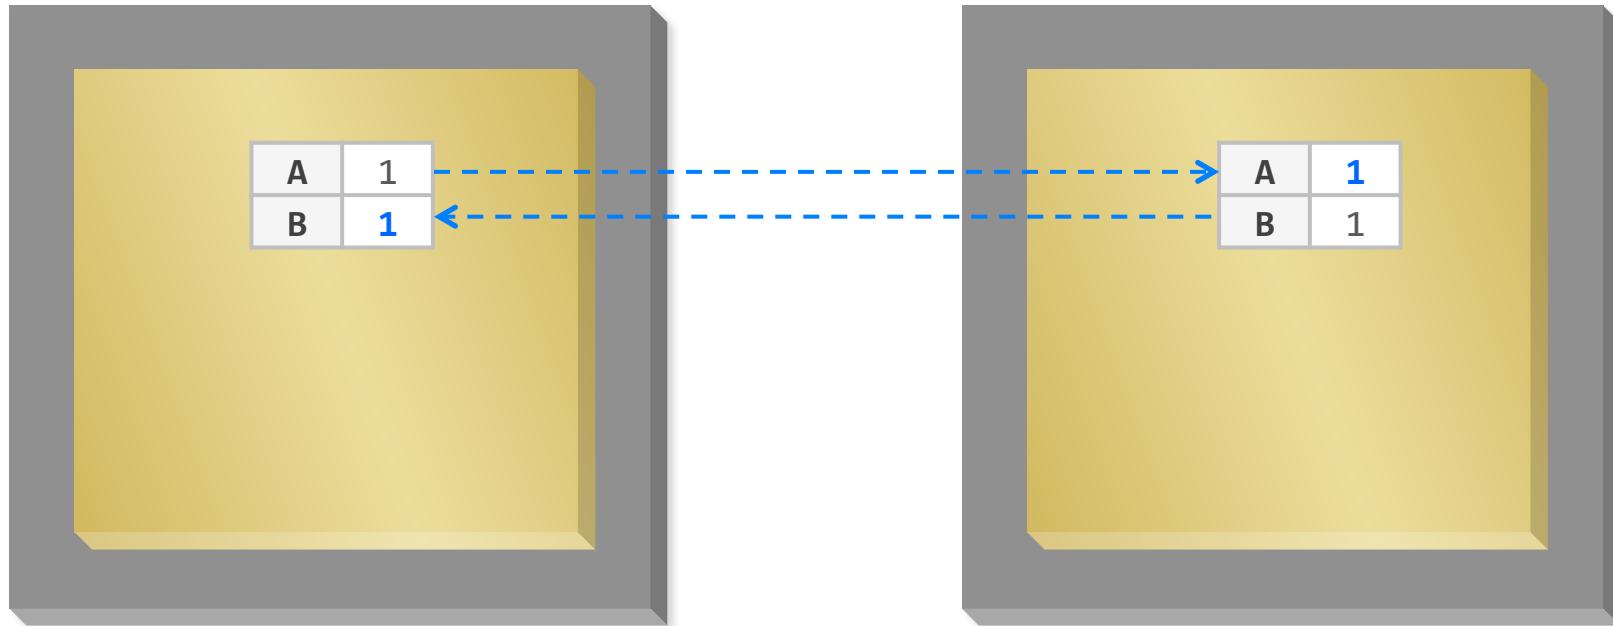
# VISUALIZING LOW-LEVEL ATOMICS

This analogy corresponds to each CPU core having its own cache.



# VISUALIZING LOW-LEVEL ATOMICS

Eventually, changes propagate between threads, but the timing is unpredictable.



# SEQUENTIALLY CONSISTENT ATOMICS

*The magic compilers use to implement them*

	Load	Store
x86/64	mov %, %	lock xchg %, %
PowerPC	hwsync ld %, % cmp %, 0 bc # isync	hwsync st %, %
ARMv7	ldr %, % dmb	dmb str %, % dmb
ARMv8	ldar %, %	stlr %, %
Itanium	ld.acq %, %	st.rel %, % mf

# HOW TO CONVERT GAME ATOMICS

*to low-level C++11 atomics*

## Game Atomics

```
volatile int A;  
  
A = 1;  
int a = A;  
  
LIGHTWEIGHT_FENCE();  
FULL_FENCE();  
  
AtomicIncrement(A);  
AtomicCompareExchange(A, ..., ...);  
...
```

## Low-Level C++11 Atomics

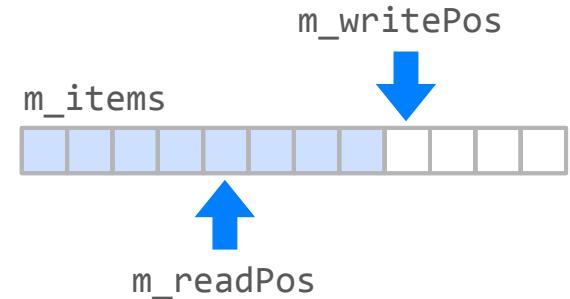
```
atomic<int> A;  
  
A.store(1, memory_order_relaxed);  
int a = A.load(memory_order_relaxed);  
  
atomic_thread_fence(memory_order_acquire/release);  
atomic_thread_fence(memory_order_seq_cst);  
  
A.fetch_add(1, memory_order_relaxed);  
A.compare_exchange_strong(..., ..., memory_order_relaxed);  
...
```

# EXAMPLE

## Capped wait-free queue in C++11

```
template <class T, int size>
class CappedSPSCQueue
{
private:
    T m_items[size];
    atomic<int> m_writePos;
    int m_readPos;

public:
    CappedSPSCQueue() : m_writePos(0), m_readPos(0) {}
    bool tryPush(const T& item) { ... }
    bool tryPop(T& item) { ... }
};
```



# EXAMPLE

## Low-level with standalone fences

```
bool tryPush(const T& item)
{
    int w = m_writePos.load(memory_order_relaxed);
    if (w >= size)
        return false;
    m_items[w] = item;
    atomic_thread_fence(memory_order_release);
    m_writePos.store(w + 1, memory_order_relaxed);
    return true;
}
```

Could even be non-atomic

```
bool tryPop(T& item)
{
    int w = m_writePos.load(memory_order_relaxed);
    if (m_readPos >= w)
        return false;
    atomic_thread_fence(memory_order_acquire);
    item = m_items[m_readPos];
    m_readPos++;
    return true;
}
```

# EXAMPLE

## Low-level with standalone fences

```
bool tryPush(const T& item)
{
    int w = m_writePos.load(memory_order_relaxed);
    if (w >= size)
        return false;
    m_items[w] = item;
    atomic_thread_fence(memory_order_release);
    m_writePos.store(w + 1, memory_order_relaxed);
    return true;
}
```

When **this** load sees the value written by **this** store...

... the fences synchronize with each other (§29.8.2, N3337).

```
bool tryPop(T& item)
{
    int w = m_writePos.load(memory_order_relaxed);
    if (m_readPos >= w)
        return false;
    atomic_thread_fence(memory_order_acquire);
    item = m_items[m_readPos];
    m_readPos++;
    return true;
}
```

# EXAMPLE

## Low-level ordering constraints

```
bool tryPush(const T& item)
{
    int w = m_writePos.load(memory_order_relaxed);
    if (w >= size)
        return false;
    m_items[w] = item;
    m_writePos.store(w + 1, memory_order_release);
    return true;
}
```

When the load sees the value written by the store...

... the store synchronizes-  
with the load (§29.3.2).

```
bool tryPop(T& item)
{
    int w = m_writePos.load(memory_order_acquire);
    if (m_readPos >= w)
        return false;
    item = m_items[m_readPos];
    m_readPos++;
    return true;
}
```

# EXAMPLE

## Using sequentially consistent atomics

```
bool tryPush(const T& item)
{
    int w = m_writePos;
    if (w >= size)
        return false;
    m_items[w] = item;
    m_writePos = w + 1;
    return true;
}
```

When the load reads from the store, they synchronize-with each other (§29.3.1).



```
bool tryPop(T& item)
{
    int w = m_writePos;
    if (m_readPos >= w)
        return false;
    item = m_items[m_readPos];
    m_readPos++;
    return true;
}
```

# EXAMPLE

## *Capped wait-free queue in C++11*

```
template <class T, int size>
class CappedSPSCQueue
{
private:
    T m_items[size];
    atomic<int> m_writePos;
    alignas(64) int m_readPos;

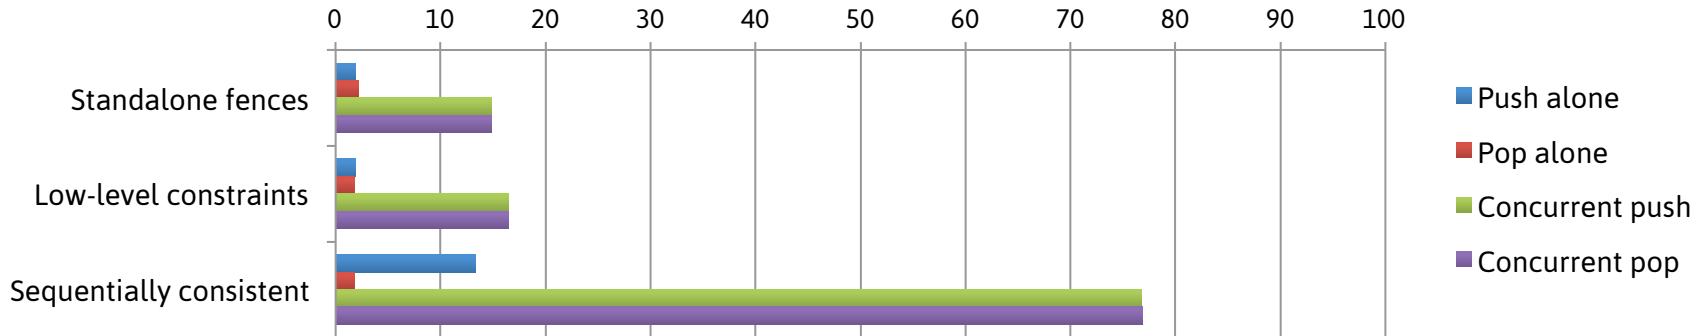
public:
    CappedSPSCQueue() : m_writePos(0), m_readPos(0) {}
    bool tryPush(const T& item) { ... }
    bool tryPop(T& item) { ... }
};
```

All other variables can remain non-atomic because there is no data race.

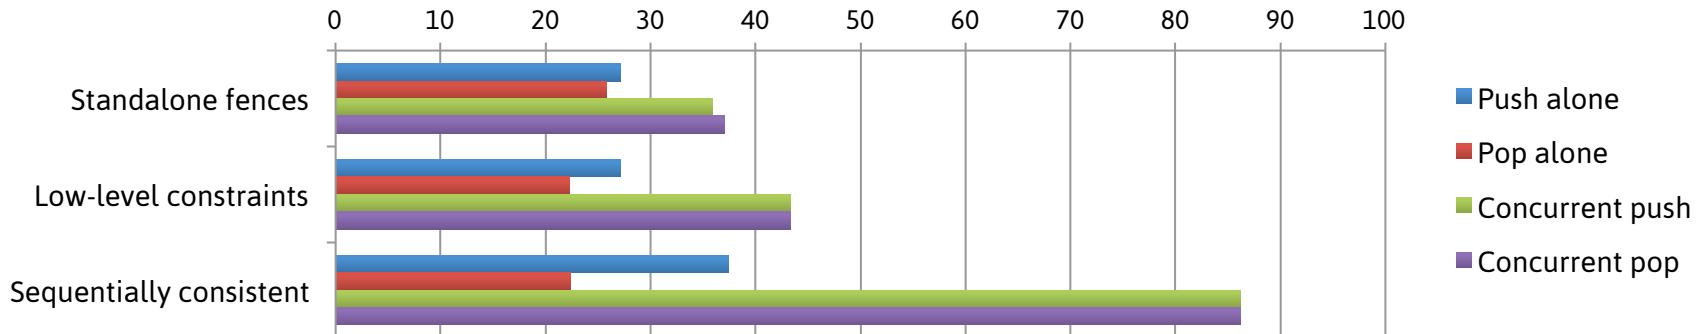
# BENCHMARKS

*nanoseconds per operation*

**Intel Core-i7** Quad-core / 2.3 GHz / Xcode 5.1.1 Release



**ARM Cortex-A9** Dual-core / 800 MHz / Xcode 5.1.1 Release



# RECAP: The C++11 Atomic Library

- C++11 forbids “data races”
- Two atomic libraries
- Pass non-atomic information by *synchronizing-with*

# THANKS



**UBISOFT®**

Hugo Allaire  
Dominic Couture  
Jean-François Dubé  
Dominique Duvivier  
Michael Lavaire  
Jean-Sébastien Pelletier  
Rémi Quenin  
James Therien

Charles Bloom  
Hans Boehm  
Bruce Dawson  
Peter Dimov  
Maurice Herlihy  
Paul McKenney  
Peter Sewell  
Herb Sutter  
Anthony Williams  
Dmitry Vyukov

**Jeff Preshing**

 @preshing

 jeff@preshing.com

**Preshing on Programming**



<http://preshing.com/>